

Intelligent Clustering in Wireless Sensor Networks

by

Robert William Lee Guderian

A thesis submitted to
The Faculty of Graduate Studies of
The University of Manitoba
in partial fulfillment of the requirements
of the degree of

Master of Science

Department of Computer Science
The University of Manitoba

Winnipeg, Manitoba, Canada

July 2012

© Copyright by Robert William Lee Guderian, 2012

Thesis advisor

Dr. Rasit Eskicioglu

Author

Robert William Lee Guderian

Intelligent Clustering in Wireless Sensor Networks

Abstract

Wireless Sensor Networks (WSNs) are networks of small devices, called motes, designed to monitor resources and report to a server. Motes are battery-powered and have very little memory to store data. To conserve power, the motes usually form clusters to coordinate their activities. In heterogeneous WSNs, the motes have different resources available to them. For example, some motes might have more powerful radios, or larger power supplies. By exploiting heterogeneity within a WSN can allow the network to stay active for longer periods of time.

In WSNs, the communications between motes draw the most power. By choosing better clusterheads in the clusters to control and route messages, all motes in the network will have longer lifespans. By leveraging heterogeneity to select better clusterheads, I have developed Heterogeneous Clustering Control Protocol (HCCP). HCCP is designed to be highly robust to change and to fully utilize the resources that are currently available.

Contents

Abstract	ii
Table of Contents	iv
List of Figures	v
List of Tables	viii
Acknowledgments	ix
Dedication	x
1 Introduction	1
2 Related Work	7
2.1 Wireless Sensor Networks	8
2.2 MAC Layer Protocols	9
2.3 Non-clustering Protocols	11
2.4 Clustering Protocols and Communication	14
2.5 Heterogeneous Wireless Sensor Networks	16
3 HCCP: Heterogeneous Clustering and Control Protocol	19
3.1 Theory of Operation	21
3.2 Implementation Details	28
3.2.1 Clusterhead Election	28
3.2.2 HCCP Goodness Delay	31
3.2.3 Roundtable Discussion Functions	36
3.2.4 Clusterhead Choice Timeout	38
4 Experimental Setup and Results	40
4.1 Simulation of HCCP and LEACH	41
4.1.1 Modifications to LEACH for Simulation	46
4.1.2 Modifications to HCCP for Simulation	47
4.1.3 Observations and Results	47
4.1.4 Discovering how Heterogeneous Factors Effect the Network . .	51
4.1.5 Tuning HCCP for Power Efficiency	62

4.2	Testbed Setup, Deployment and Results	74
4.2.1	Building HCCP with 802.15.4	74
4.2.2	Description of Testbed Motes	78
4.2.3	Heterogeneous Setup of the Motes	86
4.2.4	Testbed deployment	87
5	Conclusion and Future Work	95
5.1	Future Goals of WSNs	98
	Bibliography	100

List of Figures

1.1	High-level description of a mote. Adapted from [1].	2
1.2	A hierarchical cluster topology [2].	3
2.1	Device A and device B can both communicate with device C, but A and B are out of range of each other.	9
2.2	A CSMA handshake between two devices.	10
2.3	A sample of a TDMA schedule	10
3.1	HCCP uses existing MAC protocols and can be used by any application, acting as middleware.	20
3.2	LEACH election run cycle.	22
3.3	A visualization of the HCCP election and run cycle.	24
3.4	An invisible wall created due to HCCP clusterhead candidacy messages.	29
3.5	HCCP messages received at the base station. There is no way messages outside the radio range of the sink can send messages to the sink.	30
3.6	LEACH messages reached at the base station, note there is no ‘wall’ blocking messages.	31
3.7	M3 never becomes a clusterhead since M1 and M2 announce their Clusterhead Candidacy first. Since M3 is never a clusterhead, M4 never has a chance to send its messages through the network.	34
3.8	A mote range 300 away would take at best 3 rounds to receive a beacon given a range of 100.	37
3.9	Mote C would choose mote D as a clusterhead even though mote B is closer to the sink.	39
4.1	Visualization showing the relationship between Percentage of Clusterheads per round and simulation results.	49
4.2	The Goodness Delay focuses the clusterhead task on motes that are suited for the task.	50
4.3	The relationship between HCCP Goodness Delay using Available Queue size and simulation results.	52

4.4	HCCP handles being overloaded with messages well, as motes with larger available message queues will opt to be clusterheads.	53
4.5	LEACH does not handle being overloaded, as it does not consider message queue size when electing clusterheads. Note the y axis is the same as Figure 4.4.	53
4.6	HCCP loses fewer messages due to full message queues than LEACH.	54
4.7	The relationship between HCCP Goodness Delay using Available Battery Power size and simulation results.	55
4.8	HCCP focused on Battery Power loses less messages since dying motes will not elect to be clusterheads.	57
4.9	HCCP focused on Battery Power has better throughput than LEACH since fewer messages are lost in dying motes.	57
4.10	Visualization showing the relationship between HCCP Goodness Delay using Duty Cycling and simulation results.	58
4.11	The relationship between HCCP Goodness Delay using Random Wait Times and simulation results.	60
4.12	The relationship between HCCP Goodness Delay using Sensor Mission and simulation results.	61
4.13	Tuning the efficiency of HCCP by minimizing Roundtable discussion allows HCCP to have networks lifespans approximately equal to LEACH.	63
4.14	In a tuned network, HCCP still has a much higher message throughput.	63
4.15	HCCP using Clusterhead Candidacy vs Not using Clusterhead Candidacy. The message throughput is comparable between the two networks.	66
4.16	HCCP using Clusterhead Candidacy vs Not using Clusterhead Candidacy. The results are quite even in terms of mote lifespan.	66
4.17	The percentage of messages created that were received at the sink is approximately the same whether the sink periodically sleeps or not.	69
4.18	Average mote lifespan is unaffected by whether the sink periodically sleeps or not.	70
4.19	Snapshot of a part of a WSN: routing mote R1 can not communicate with routing mote R2 because the routing backbone is broken.	72
4.20	The Sun Microsystems SunSPOT.	79
4.21	A SunSPOT free-range mote and SunSPOT base station.	79
4.22	A SunSPOT taken apart to view the components in the casing.	80
4.23	Seeeduino Stalker with XBee port exposed.	81
4.24	An Arduino Duemilanove turned into a mote using the XBee Shield and an XBee radio.	83
4.25	The parts used to make the Arduino mote.	84
4.26	The tmote Sky, which is a Telos variant.	86
4.27	Motes distributed around a house.	89
4.28	Cumulative number of messages received from the deployment done in a house.	90

4.29 The grid deployment setup.	92
4.30 Results from LEACH and HCCP Grid Deployments	93

List of Tables

4.1	Breakdown of data packets in 802.15.4. Adapted from [3].	75
4.2	Customized 802.15.4 packets for use in HCCP with more detail. . . .	75
4.3	Sample TDMA schedule as sent by the clusterhead to its clustermotes.	77
4.4	Comparison of Motes in Testbed Deployment	88

Acknowledgments

I would like to begin by thanking my Dr. Rasit Eskicioglu who has been my long-suffering advisor to whom I will forever be grateful. I'd like to thank my committee, Dr. Peter Graham and Dr. Bob McLeod for all their input and patience. I'd finally like to thank my biggest fan, my inspiration, my editor, my best friend and wife, Becky.

Dedicated to Becky for always believing in me.

Chapter 1

Introduction

A Wireless Sensor Network (WSN) is a network of many small, battery-powered devices, called motes, that create ad-hoc networks via wireless radios. Motes are inexpensive devices outfitted with various sensors and are configured to monitor some phenomena and report back to a server – called a sink. Estrin et al. [4] elegantly described WSNs as,

“...sensors only interact with other sensors in a restricted vicinity, but nevertheless collectively achieve a desired global objective.”

This captures the essence of a WSN concisely, describing the multi-hop environment and how a WSN consists of many small objects that are working towards a larger goal.

Motes are simple devices that have very few components. Motes consist of a battery for power, sensors for detecting some phenomena, memory for storing the data collected, a radio for communication and a processor for controlling all the other components. Optionally motes can have ways to regenerate power, such as a solar

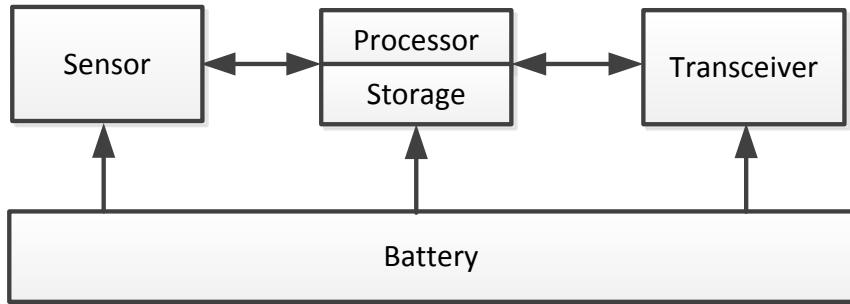


Figure 1.1: High-level description of a mote. Adapted from [1].

panel, and some motes have the ability to move. The layout of a mote can be seen graphically in Figure 1.1.

Applications for WSNs range from animal habitat monitoring [5, 6] to building and bridge structural monitoring [7], to monitoring live volcanoes during eruptions [8]. Typically, when deployed, a WSN self-configures, creating routes from sensor motes to the sink automatically. Most routes in WSNs are not direct to the sink, therefore motes must use several peer motes as a multi-hop network to send messages to the sink. Routing protocols in WSNs must automatically create these multi-hop routes to form a functioning network. To create these routes, routing algorithms designed for WSNs, such as the Collection Tree Protocol [9] are used.

There are many scenarios where a WSN could be used to make a task simpler, or even save lives. For instance, Werner-Allen et al. [8] have used a WSN to monitor a seismic and infrasonic (low-frequency acoustic) volcanic eruption. The team wanted to have as many motes as possible, and as long a network lifespan to collect as much data as possible. Every project has a limited budget, and high-powered radios are expensive and use lots of energy. Since high-powered radios were ruled out, Werner-Allen et al. chose to use a multi-hop WSN. A WSN allowed the team to deploy

many sensors with low-power radios that could communicate their findings back to the computer they used as a sink. Some sensors were lost during the eruption, which could have meant a loss of life or a loss of data if the sensors required people to continually check them. Since a WSN was used, the data was transferred off the mote before it was lost to the eruption and provided some very interesting data!

WSNs can also be used in less exciting work, such as bridge or vehicle monitoring [7]. If it is difficult or inconvenient for a person to check strain levels on bridges, a WSN can be deployed to monitor the bridge. Some of these bridges can be difficult to get power to, as they might be far into a forest or in a remote location. A battery-powered WSN does not require external power to be run to the area, and can continually monitor strain levels without human intervention, saving companies or governments time and money while making the structure safer.

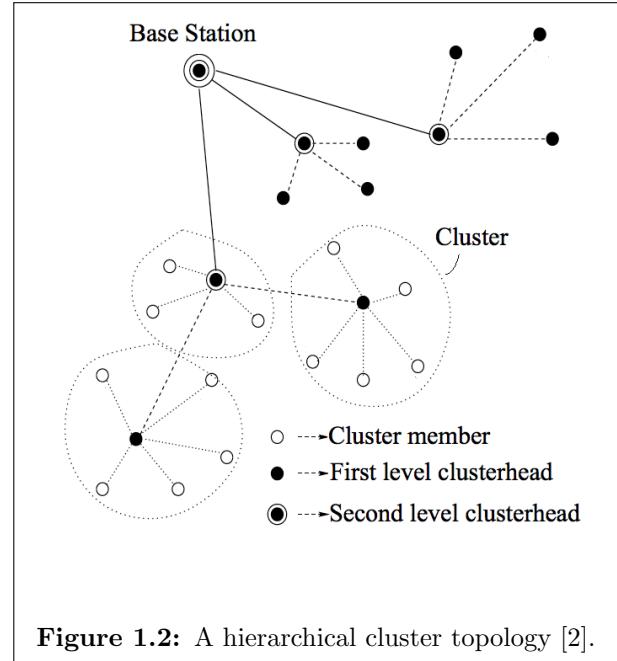


Figure 1.2: A hierarchical cluster topology [2].

To simplify routing and to conserve energy, a WSN groups motes into clusters. These clusters elect a clusterhead which works as a gateway and routes all network traffic in and out of the cluster. Motes that are cluster members only communicate with the clusterhead. Clusterheads then route the messages through other cluster-

heads, to the sink. An example of this can be seen in Figure 1.2. This simplifies the network topology as each cluster can be viewed as one large mote. Messages to the sink are then routed through clusterheads only, which lowers the number of hops taken from the originating cluster to the sink. Using clusterheads reduces the energy draw from all motes that are in the cluster since they only need to send messages to the clusterhead. The clusterhead accepts messages from the motes in the cluster, creates one large message out of all the messages, and sends this large message to the sink. Handling messages in this way allows motes that are cluster members to sleep or be idle while the clusterhead attempts to send the message across the network, thus saving power.

A heterogeneous WSN has motes that are not all identical. There may be several different types of motes from different companies or motes from the same company with different characteristics. Motes may be running different operating systems, be programmed with different software or have additional hardware. Homogeneous networks can also be viewed as heterogeneous networks; heterogeneity occurs as motes change over time, such as a mote's battery depletion or a mote malfunctioning. These differences, large and small, can be exploited to extend the network lifespan or increase message throughput.

In a WSN, typically a mote is assigned to provide detailed information about the phenomenon it is monitoring. If a few motes are equipped with more expensive hardware, the network should adapt to allow these motes to live longer. This can be done by not allowing the more expensive mote to be a clusterhead, and by providing a clusterhead in range of the expensive mote. As this specialization increases, so

does the heterogeneity of the network. Clusterheads should be selected to allow the specialized motes to have longer lifespans. Or, if a mote has a larger available message queue it should be elected to be a clusterhead since it will be less likely to lose messages if a large number of messages are sent to it, making the mote a good message routing mote.

A ‘long-lived WSN’ is a network that is designed to stay active for extended periods of time. Long-lived WSNs use techniques such as duty cycling and make as few transmissions as possible to achieve longer network lifespan. It is difficult to have a long-lived network mostly due to the battery-powered nature of motes. Long-lived networks often trade off some functionality, such as the frequency of sampling with their sensors and the frequency of reporting, to increase mote lifespan.

A clusterhead is a mote that controls transmissions from other motes in the immediate neighbourhood. Clusterhead motes draw more power due to the increased number of transmissions they perform as they must stay active to receive messages from the cluster. The key to long-lived networks is selecting motes that are well-suited to be clusterheads. Therefore, in a heterogeneous long-lived WSN, motes that have more or larger batteries and have plenty of residual energy should be elected to be clusterheads. If the clusterhead has little power left or has a malfunctioning radio, the probability of losing messages increases. Using heterogeneity to avoid choosing a mote that is better designed to be a clusterhead will aid in designing a long-lived network.

Suitable clusterheads are important when building a long-lived WSN. The clusterhead draws a large amount of power, therefore the performance of clusterhead motes

will degrade over time. If a clusterhead mote detects that it is performing poorly (e.g., cannot keep up handling the incoming messages from their motes), it should opt-out or demote itself to pass the task on to a more capable mote. Passing the clusterhead task to a more capable mote will prevent the poorly-performing mote from losing transmissions and causing retransmissions, which creates unnecessary overhead in the surrounding motes.

To increase message throughput, extend lifespans of individual motes and allow WSNs to function longer, I have developed Heterogeneous Clustering and Control Protocol (HCCP). HCCP uses a taxonomy to describe heterogeneity and selects candidate clusterheads based on the capabilities of motes.

Chapter 2

Related Work

Many researchers have contributed to the WSN field, all focusing on different areas. There are a few different aspects of WSNs that are considered in this thesis. To reflect this, the related work is grouped by different areas of WSN research.

Areas of WSN research include radio transceiver technology which has been developed for decades, but has only recently been developed as a low-power technology. WSN research also includes MAC layer protocols, which are used by radios to share a radio channel between multiple devices. To move messages towards the sink, WSN nodes can either create clusters (Clustering Protocols), or communicate without clusters (Non-clustering Protocols). There are benefits to using clusters in WSNs, and other benefits to not using clusters. Finally, heterogeneity in WSNs can be used to improve the performance of a WSN.

2.1 Wireless Sensor Networks

Wireless technologies are by no means a new invention. In 1892 Nikola Tesla proposed that radio waves could be used for communication without any wires connecting the two points [10]. Since then, wireless radios have evolved from large, power hungry devices such as radio backpacks used in World War II, to the ubiquitous cell phone of today. Wireless radios have gotten smaller, more power efficient and much more affordable.

The story of the computer has much in common with the evolution of the wireless radio. Starting with computers that were large and extremely expensive, computers are now standard equipment for everyday use. The cost of computers has gone way down, as has the size of computers. Both wireless radios and computer are core components of WSN nodes.

The paper that is often credited in being the first paper on designing WSNs is Pottie and Kaiser's [11] *Wireless Integrated Network Sensors*. Pottie and Kaiser proposed the acronym WINS for the field, and were the first to connect the ideas of pervasive low power computing with sensor networks and proposed an architecture to help solve these problems. At its core WSNs are just microcontrollers, sensors and wireless radios; Pottie and Kaiser saw the potential of these objects, and showed the world that the combination is more than the sum of its parts. The authors moved beyond looking at the components, and considered what could be possible with the concept, and how it would need to be achieved. Pottie and Kaiser saw that the density of the distributed network is the strength of WSNs. Key points of the paper were that using the shortest transmission range is important to increase node lifespan, and

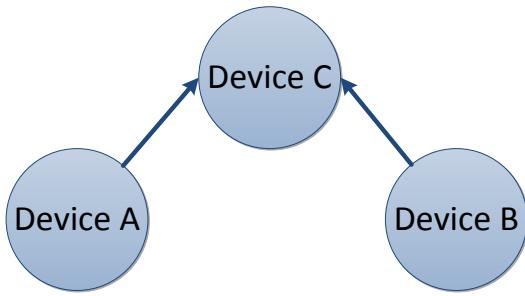


Figure 2.1: Device A and device B can both communicate with device C, but A and B are out of range of each other.

that networks must be self-organizing to work efficiently.

2.2 MAC Layer Protocols

Carrier sense multiple access (CSMA) [12, 13] is a general protocol that is general to networking devices that have a shared channel. Designed by Kleinrock and Tobagi in 1975 for transmitting data over radio, CSMA requires devices to listen on the channel to ensure no other devices are transmitting before beginning a transmission. The ‘carrier sense’ part of the name comes from the device that is ready to send checking (or sensing) to see if there is already a transmission in progress. Since the protocol is designed to have many devices sharing the same channel, it is a ‘multiple access’ protocol. CSMA has a problem called ‘the hidden node problem’. This arises when two devices that are out of radio range of each other are communicating with the same device that is in between the two. This problem can be seen in Figure 2.1, where device A and B are both trying to communicate with device C. Device A and B do channel checks and do not hear any noise on the channel, so they begin transmitting. Device C gets both transmissions at the same time, causing a collision

in the messages, and neither of the messages are received.

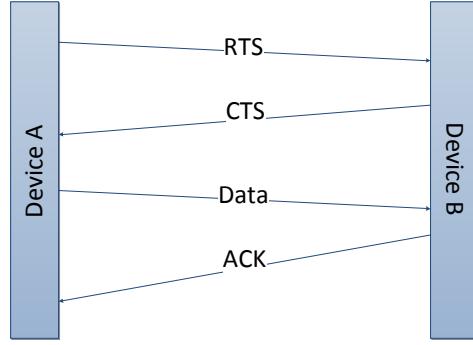


Figure 2.2: A CSMA handshake between two devices.

CSMA has been adapted to WSNs by Woo and Culler [14], adding some enhancements to make transmissions more reliable. Woo and Culler added the idea of a two-way handshake to transmit messages across the network. The device sends a ‘ready to send’ (RTS) message to the intended recipient (which is heard by all surrounding devices due to the shared medium). If the intended recipient is not currently busy, it sends back a ‘clear to send’ (CTS) message. CSMA handshaking is illustrated in Figure 2.2. Woo and Culler also show that in WSNs that are largely collision-free, the ACK step of the handshake can be dropped for energy savings. CSMA remains a major part of most WSN communication due to its simplicity and effectiveness in detecting collisions.

\leftarrow Follow schedule \rightarrow						\leftarrow Follow schedule \rightarrow					
TDMA Schedule Announcement						TDMA Schedule Announcement					
padding	Sensor 1	Sensor 2	Sensor 3	Sensor 4	...	padding	Sensor 1	Sensor 2	Sensor 3	Sensor 4	...

Figure 2.3: A sample of a TDMA schedule

Time division multiple access (TDMA) [15] is another method for allowing mul-

tiple devices to share a radio channel. Since the channel is shared, only one device can successfully transmit at a time. If more than one device is transmitting at the same time, the messages will collide, and neither message will be received correctly. To avoid having two devices transmitting at the same time, a coordinator device announces a schedule that defines when and how long the devices in the immediate area are allowed to transmit. A visual representation of TDMA can be seen in Figure 2.3.

2.3 Non-clustering Protocols

Sparse Topology and Energy Management (STEM), designed by Schurgers et al. [16] is a simple MAC protocol that uses peer to peer messaging without a clusterhead to send messages across a network. Since having the radio on drains the battery quickly, all motes cycle between listening and sleeping to conserve energy. The motes do not synchronize the schedule for sleeping and listening times, so motes poll neighbouring motes to send messages across the network. Polling wastes transmissions, since a mote will continue to send messages to a mote that is asleep until the mote wakes up and responds to the polling. No synchronization means that energy is wasted sending messages that are not received by anything.

Created by Intanagonwiwat et al. [17, 4], Directed Diffusion uses peer-to-peer communication to transmit messages across the network. Directed Diffusion does not use clusterheads or coordinator devices to collect or route information. Since there are no clusterheads, all the devices in the network are tasked with tracking where the sink is in the network, and then transmit messages to devices in the direction that the sink is. Intanagonwiwat et al. use a type of beaconing as a simple routing protocol. The

sink announces its position periodically to its neighbouring nodes. The neighbouring nodes then announce to their neighbours that they are one hop away from the sink. This pattern continues until the entire network knows how many hops away from the sink they are. Directed Diffusion provides effective routing information with low overhead. Intanagonwiwat et al. considered sending some extra information with the beacons, such as which sinks are interested in what information being sensed, but did not consider sending extra information about the motes along with the beacon.

Gossiping is a simple way of increasing the chances that a message will be received at the sink. Hedetniemi et al. [18] discussed the simplicity and simple gains achieved by using gossiping in their survey paper on gossiping. Gossiping, in general, is sending the same message along multiple routes to the sink. In a network with packet loss, sending multiple copies of the same message will increase the odds the message will be received at the sink. This also has the side-effect of possibly having the same message received at the sink multiple times. Another negative side-effect is that more transmissions, and therefore more power is needed to send one message to the sink. Though an effective way of sending messages, the cost involved in sending the messages must be considered before using a gossiping technique.

To quickly disseminate data, flooding [1] can be employed to broadcast a message across an entire WSN. Flooding can be considered the extreme case of gossiping, where a message is announced to all surrounding devices, which in turn announce the same message. This is a good method for sharing routing information quickly, but is expensive. All the devices in the network must be on, and will receive and transmit the message. Flooding also tends to cause many collisions in the network, as all the

devices will repeat the message shortly after receiving the message.

S-MAC [19] is a reliable way of moving data across a WSN without using clusterheads. Ye et al outline the important features of a successful MAC layer protocol: energy efficiency, fairness (all motes have an equal chance of messages reaching the sink), low message latency and high throughput. Due to the constrained nature of WSNs, it is difficult or impossible to have all of these features, though networks can be tuned to enhance the desired features the network should have. An example of this is that a network could have better throughput, but lowered energy efficiency. S-MAC does sleep cycling to improve the lifespan of the network, with periodic listening states to check if any devices are trying to send to it. To avoid polling, S-MAC uses synchronization techniques to wake all the devices at approximately the same time, so no devices will be polling a device that is sleeping.

Power Efficient Gathering in Sensor Information Systems (PEGASIS), designed by Lindsey and Raghavendra [20], uses global knowledge of the network to create paths to the sink. PEGASIS saves energy by using greedy algorithms to create a near-optimal path to the sink. Motes hop messages to neighbours that are closer to the sink. When a mote receives a message, it performs message fusion on the message, then sends the fused messages to a neighbour that is closer to the sink. Motes using PEGASIS accept many messages and send one message per round due to this data fusion and hopping setup. Reducing the number of messages sent reduces the amount of energy used while the network runs. PEGASIS' use of global knowledge is unrealistic in most network deployments (e.g. WSN nodes launched from a plane, thrown into a bush) and therefore limited in its real-world applications.

2.4 Clustering Protocols and Communication

The standard approach for clustering, designed by Heinzelman et al. [21], is Low-Energy Adaptive Clustering Hierarchy (LEACH). LEACH is a commonly used algorithm for clustering and low energy communication. Most other clustering algorithms are in some way extensions of LEACH. Being the earliest clustering protocol for WSNs, LEACH uses somewhat simplistic methods for electing clusterheads. This simplicity is actually the key strength of LEACH, since the clusterhead election makes no assumptions about the network. Heinzelman et al. did not consider heterogeneity or mote capabilities for LEACH.

Some protocols focus on sensing and reporting anomalies. Anomalies sensed could be fires, earthquakes, security alerts, etc. Threshold sensitive Energy Efficient sensor Network (TEEN) [2], developed by Manjeshar and Agrawal, focuses on only reporting anomalies that are worth reporting, not wasting transmissions on data that is not of interest. Since staying on and transmitting messages uses the most energy in WSNs, TEEN only sends messages if the sensor readings are important enough to send. TEEN uses a threshold to define whether or not a sensor reading is valuable and worth sending to the sink or not. TEEN is a reactive-style WSN, it is not designed to regularly send information to the sink, only if an anomaly that is being monitored occurs will a message be sent to the sink. This is useful for monitoring for fires or intruders and the like, but not useful for monitoring a resource over time.

Soro and Heinzelman [22] reported that, when only using residual energy in a mote for electing a clusterhead, the lifespan of the network is negatively affected when compared to using more than one factor. This demonstrated that the benefits of

using hybrid criteria when choosing a clusterhead outweigh the overhead it generates. While Soro and Heinzelman showed that comparing more than one factor between motes in an election generates excessive overhead, they did not consider summarizing multiple factors into one hybrid criteria.

SPIN, designed by Kulik et al. [23] uses an advertisement phase where a mote informs surrounding motes about the message that it has. The surrounding motes then can request the message from the mote that sent the advertisement. There is metadata in the advertisement that has some details about the message that can be requested. This advertisement phase is interesting in that any metadata could be in the advertisement message, such as information about the mote, or messages to surrounding motes. Kulik et al. only considered sending information about the message that is ready to be sent.

Building on the foundation of LEACH, Younis and Fahmy [24] created Hybrid Energy-Efficient Distributed clustering (HEED) to address the issue of selecting better clusterheads. HEED uses a hierarchy much like LEACH, but uses more intelligence to choose the next clusterhead. HEED uses residual power and a secondary parameter to create a single value describing how well suited the mote is to being a clusterhead. HEED has addressed selecting better clusterheads based on certain parameters, but only in homogeneous networks.

Dong and Liu [25] created a model that chooses clusterheads based not only on a mote's capacity to be a clusterhead, but on data it has collected while the network has been alive. If a mote had previously been chosen as a clusterhead, but did not do well, that mote goes to a blacklist and will only be chosen again if there are no

known candidates to be a clusterhead. This use of historical data allows the network to improve its clusterheads over time, creating a network of clusterheads that are known to work well. Though Dong and Liu used knowledge of motes in the network to improve network lifespan, they did not consider heterogeneous networks where motes may have special features, such as solar panels, that could restore a mote with dead batteries to a useful state.

2.5 Heterogeneous Wireless Sensor Networks

Ou et al. [26] extended the lifespan of a heterogeneous WSN by making the network power-aware. Motes with larger power supplies were assigned the task of a clusterhead, since a clusterhead requires more energy. These clusterheads were distributed around the network to create a spanning tree. The power supply variance was the only heterogeneous aspect of the experiment and clusterheads were selected manually when the network was deployed, not dynamically by election.

Recognizing the lack of election protocols for heterogeneous WSN, Smaragdakis et al. [27] created an election protocol for heterogeneous WSNs. Extending LEACH, the protocol included remaining power levels as a weight in clusterhead elections and concluded that this weight increased network lifespan. The protocol did not consider any other heterogeneity than the residual power supply.

Brzozowski et al. [28] explored how messages should be stored in heterogeneous networks, with residual energy levels being the main focus of the heterogeneity in the network. Brzozowski et al. used a novel Media Access Control (MAC) protocol that controlled where the messages were stored and how they were routed around dying

motes. Though Brzozowski et al. looked at routing around dead motes, they did not consider using heterogenous motes to extend the network lifespan even further. Methods such as always choosing where data should be routed, or which motes should be clusterheads before devices start to have critically low energy levels.

Hu et al. [29] used two different types of motes to create a long-lived WSN. The motes were assigned different tasks in the network, some for data collection and some for routing. The authors did not consider self-configuring the network to utilize the heterogeneity or the possibility of extending the idea to more than two classes of motes.

By creating motes with differing hardware configurations, Mhatre and Rosenberg [30] showed that exploiting heterogeneity in WSNs can reduce network costs. Some of the deployed motes were inexpensive, while others were quite expensive. The expensive motes were configured to use the less expensive motes for energy-intensive tasks, allowing the more expensive motes to live longer. The focus was on the cost savings of a heterogeneous WSN, creating only a rudimentary election protocol largely based on LEACH.

Yarvis et al. [31] described three different types of heterogeneity in heterogeneous WSNs, and suggested methods for best leveraging the individual types of heterogeneity. The three types of heterogeneity identified were:

1. **Computational Heterogeneity.** Motes have different amounts of processing power available to them. Motes with more processing power should do tasks such as data fusion, or compressing messages.
2. **Link Heterogeneity.** Some motes will have more powerful radios, providing

greater transmission range. Longer range allows messages to have fewer hops in the network, which decreases the latency time between the message being created and the message arriving at its destination.

3. Energy Heterogeneity. Motes have varying amounts of energy available for their use. Some motes might be powered externally, often called ‘line’ or ‘wall’ power, while other motes might have very limited amounts of nonrenewable battery power. Motes with more battery power should take on power-intensive tasks, such as being clusterhead or performing data fusion.

Yarvis et al. focused their efforts on ways of creating optimal placement of motes with the knowledge of the heterogeneous motes in the network, and worked with S-MAC [19] as a base MAC protocol. They did not consider cluster-based WSNs as a way of further increasing the lifespan of the network.

Chapter 3

HCCP: Heterogeneous Clustering and Control Protocol

HCCP is a method for describing and advertising the internal resources of wireless sensor motes, and using the strengths of the motes to increase message throughput or extend the lifespan of the network. Once motes are aware of their internal resources, they should choose clusterheads based on the resources available. For example, motes that are assigned the clusterhead role use more energy. Therefore, motes that have extra batteries should be chosen as a clusterhead before motes without extra batteries.

A mote should not only be aware of its resources, but it should also be aware of the requirements of the network. This is a consideration because the topology of WSNs are prone to change due to the battery-powered nature of its components (e.g., a mote with a dead battery will change the topology of the WSN). Also, motes may move in the network (by some external force such as a person moving the mote), further changing the topology of the network. To handle this unpredictability, the motes

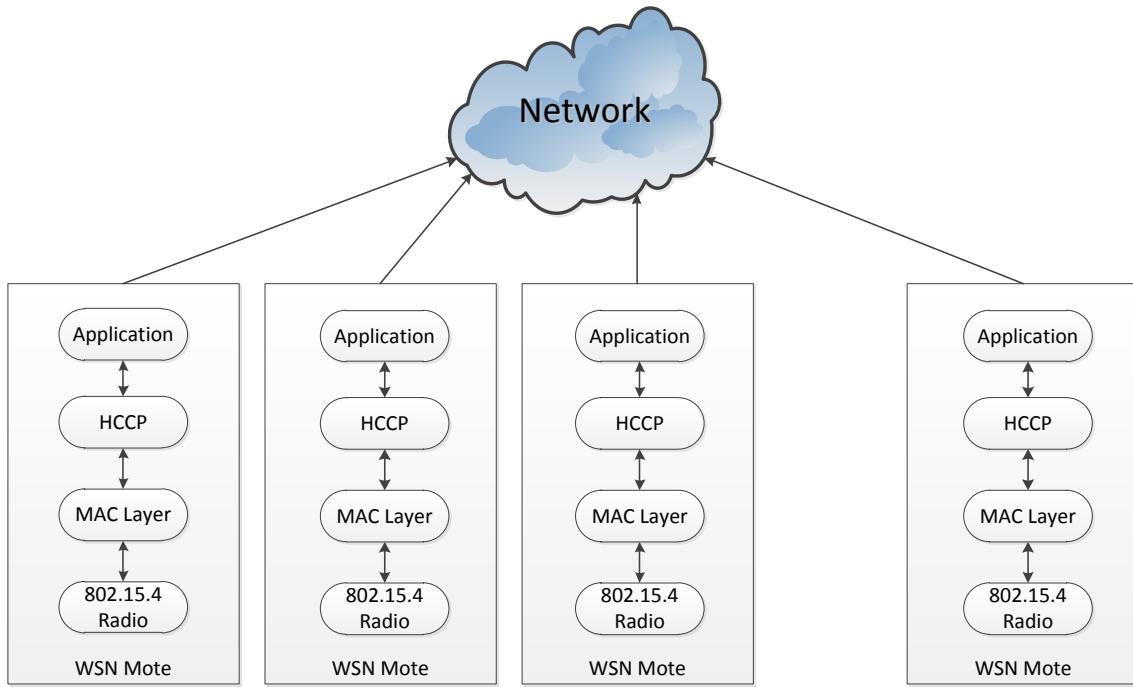


Figure 3.1: HCCP uses existing MAC protocols and can be used by any application, acting as middleware.

need to be aware of the motes nearby, but not waste excessive energy discovering them.

HCCP acts as middleware between the MAC layer and the application layer of the mote, as shown in Figure 3.1. The application layer queues messages that are ready to send, HCCP takes the messages and sends them to the sink when possible. HCCP uses existing MAC layer protocols such as CSMA and TDMA, so HCCP itself is not a MAC protocol.

3.1 Theory of Operation

HCCP has roots in LEACH, following the design of clusterhead elections, but extends the design of LEACH in multiple ways. While LEACH has relatively simple clusterhead elections, where nodes elect themselves based on a probability, HCCP has a more elaborate election process. HCCP's election process weighs the 'goodness' of the node to be a clusterhead, and automatically limits the number of other nodes that will be clusterheads. HCCP has also adopted the way a LEACH network cycles between clusterhead elections and node runtime.

The LEACH run cycle is quite simple in its operation, and provided inspiration for the HCCP run cycle. The LEACH network cycle is as follows:

1. **Clusterhead election** — Motes choose whether or not to advertise themselves as clusterheads. The decision to be clusterhead is based on how long it has been since it was a clusterhead and some added randomness.
2. **Choose cluster** — Motes that are not clusterheads send messages to the clusterhead they will follow this round.
3. **Announce Schedule** — Clusterhead sends a TDMA schedule via broadcast.
4. **Cluster Runtime** — The motes in the cluster follow the TDMA schedule, taking turns sending messages to the clusterhead.
5. **Repeat** — Go back to clusterhead election. There can be a sleep time here to extend the network life.

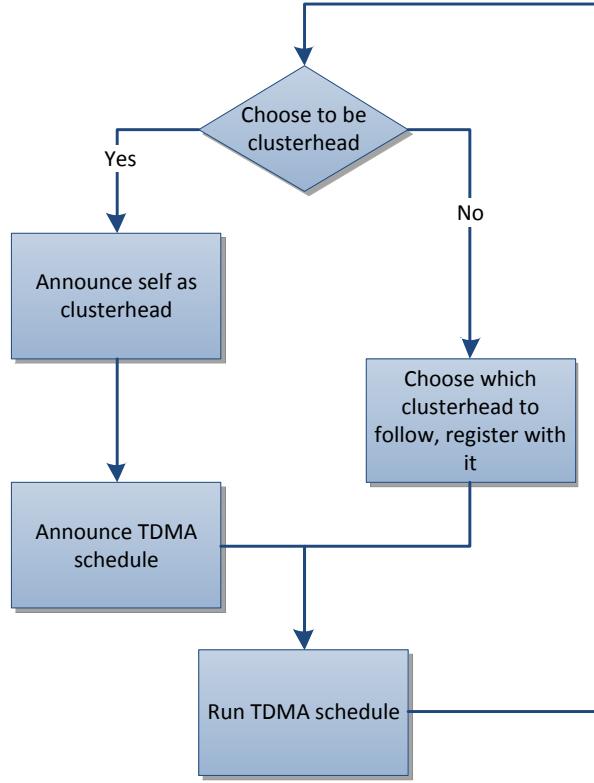


Figure 3.2: LEACH election run cycle.

THE LEACH network cycle can be seen graphically on Figure 3.2. The cycle is repeated until all the motes in the network cease to function. LEACH was focused on simplicity, which makes LEACH quite robust as it runs. Since every mote takes becomes a clusterhead occasionally, there is no single point of failure – the network will continue to function even if a number of motes cease to function. LEACH makes a solid foundation to build upon, using the lessons learned from its success and building on its weaknesses.

HCCP takes the strengths LEACH, using the simplicity of the election and run cycles, adding the ability to leverage the heterogeneity that is inherent in all WSNs. The HCCP network cycle is as follows and is illustrated in Figure 3.3.

1. Clusterhead Election

(a) **Announce Candidacy** — Announce the mote's intentions to be clusterhead. Using a Goodness assessment, announce first if the mote has a good assessment.

(b) **Announce Clusterhead** - Successful candidates announce they are clusterheads.

2. **Choose Cluster** — Same as LEACH.

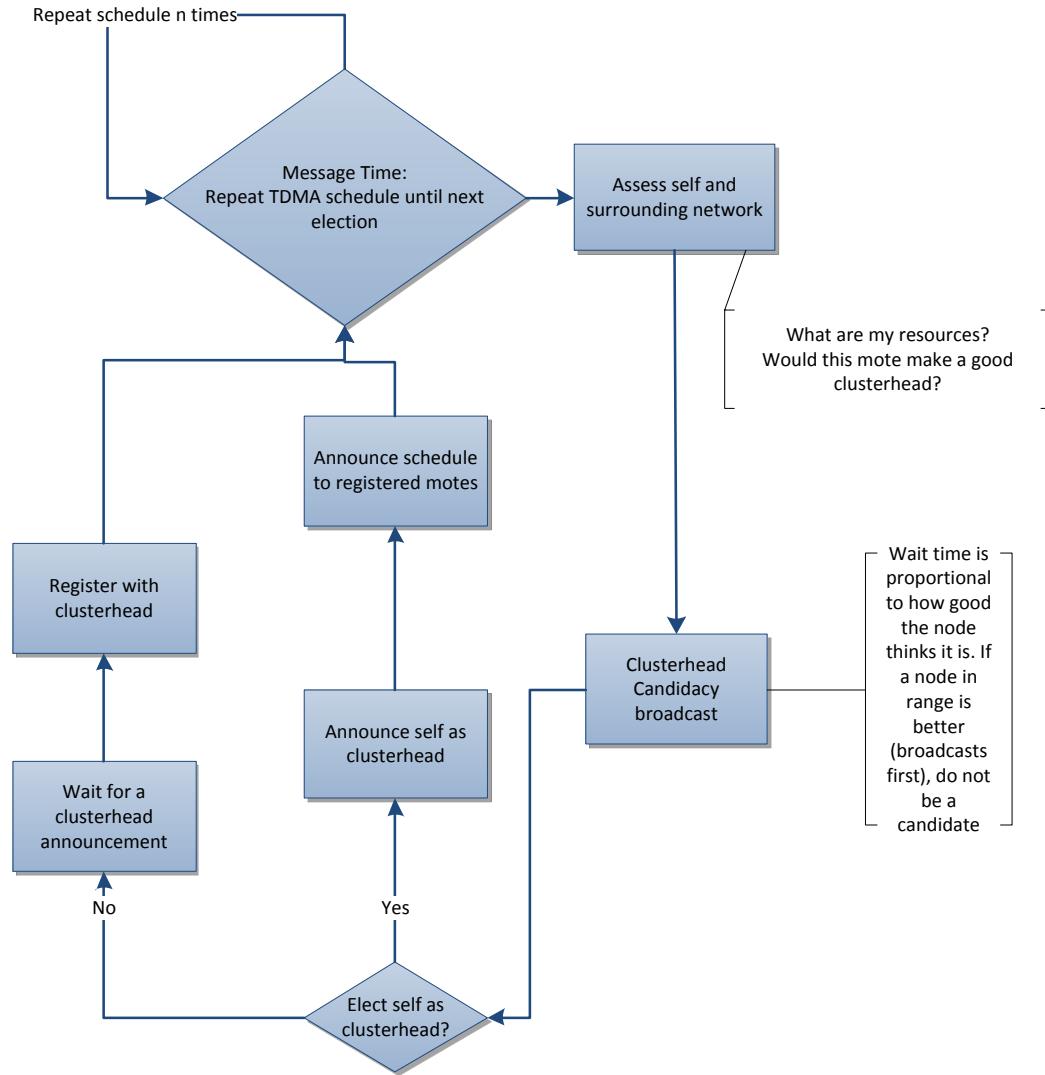
3. **Announce Schedule** — Same as LEACH.

4. **Cluster Runtime** — Same as LEACH.

5. **Roundtable Discussion** — Any queries or announcements can be made at this time. Clusterheads can opt-out, forcing a clusterhead election.

6. **Repeat** — Go back to Cluster Runtime for multiple iterations. Every n^{th} iteration (where n is selected before network deployment and well-known to the network) go back to Clusterhead election.

HCCP uses all of the strengths of LEACH in the run cycle, further modifying the concept with some extra modifications that make the network more intelligent in terms of which motes are elected clusterhead. The major additions are two-stage elections, multiple iterations of the TDMA schedule and Roundtable Discussion (including Clusterhead opt-outs).

**Figure 3.3:** A visualization of the HCCP election and run cycle.

Two-stage election

Motes choose to be a clusterhead candidate before they can elect themselves clusterheads. A mote determines that it would be a good clusterhead by inspecting its available resources, and choosing how long to wait before they announce themselves as a clusterhead. Motes that announce candidacy first get to be clusterheads that

round. Motes that determine they would not be very good clusterheads wait a longer time due to the Goodness Delay. If a candidacy announcement from a better potential clusterhead that is relatively close in proximity is overheard before a candidate broadcasts, the lesser candidate will not send an announcement, and demote itself to a regular cluster member.

Multiple iterations of schedule

All messages in clusterhead elections can be considered overhead, as these messages are not relaying any information to any destination. Further, any time a mote is on without sending or receiving messages can also be considered overhead.

To minimize time spent in clusterhead elections, they should only be run occasionally. LEACH runs a clusterhead election after every TDMA schedule runtime, which is good for distributing the burden of being a clusterhead, but creates lots of overhead messages.

The solution to this is relatively simple; a clusterhead election should only happen occasionally. Further, a TDMA schedule does not need to be re-transmitted before every TDMA runtime. A TDMA schedule only needs to be transmitted after a clusterhead election, which should only happen after n runs of the TDMA schedule. The number of times the network TDMA schedule executes (n) should be determined before network deployment. This is so it is well-known, consistent number across the entire network ensuring that the network does not get out of synchronization. A lower value of n should be chosen if the network is to be adaptive and flexible, while a larger value of n should be chosen if the network should have as little overhead as

possible.

Roundtable Discussion

After each run of the TDMA schedule, all the motes in the network turn on their radios to listen for broadcasted announcements or queries. This is the time where clusterheads can opt-out, synchronization messages can be exchanged and routing tables can be shared. The Roundtable Discussion time can also be used to extend HCCP in various ways, providing a time to implement neighbour discovery and adjust radio power accordingly, or whatever the network administrator chooses.

Clusterhead opt-out

Since the TDMA schedule published by the clusterhead will be followed multiple times, the clusterhead's performance may degrade, eventually causing the entire cluster to perform poorly. LEACH averts this problem by running a clusterhead election after every TDMA schedule runtime. Since HCCP runs multiple iterations of the TDMA schedule, the clusterhead needs a way to retire if it begins to degrade.

A clusterhead can announce its retirement during the Roundtable Discussion. When a clusterhead retires, it forces a small clusterhead election, with only the motes in the one affected cluster acting in the clusterhead election. The network does a LEACH-style clusterhead election. The first mote that announces that it is a clusterhead becomes the new clusterhead for the cluster. The Goodness Delay is used to decide how long a mote will wait before it sends a clusterhead announcement, using the same method as candidacy announcements. This way, the best possible clusterhead mote should become the replacement clusterhead.

The cluster then continues to run using the existing TDMA schedule for the remaining iterations before the next full clusterhead election.

3.2 Implementation Details

The following is a discussion of how HCCP works at an implementation level. Since HCCP is designed to be run on different hardware, with different capabilities, it is best to describe the network setup in terms of communication. Specifics about mote design, real-time operating systems and communication stacks are abstracted.

3.2.1 Clusterhead Election

HCCP takes a different approach than LEACH to electing clusterheads. Motes can't just be cluster members, even though being a clustermote every round is the most energy-efficient thing to do. A problem inadvertently created by HCCP's attempt to allow more motes be clustermembers, is that if a mote is by a sink, it would never choose to be a clusterhead. In the Clusterhead Candidacy phase, the sink will always broadcast first, never allowing motes in range of the sink to be clusterheads, since they would always concede to the sink. Since no motes around the sink would ever be clusterheads, motes out of range of the sink cannot hop messages to the sink. This would create an invisible wall around the sink, where motes past the wall will never have messages reach the sink. If a mote does not complete any messages, it is suffering starvation, a term taken from task scheduling in operating systems when processes will never run due to higher priority processes continually running. The effects of starvation can be seen in Figures 3.4 and 3.5. Figure 3.4 shows a comparison of distance from the sink to how many messages have been received at the sink, and it is clear that no messages beyond 100 arbitrary distance units from the sink would ever have messages received at the sink. Figure 3.5 shows a map of where the motes

are that have completed messages, and which motes are suffering from starvation (the sink is the green triangle in the middle of the network). A few messages from motes beyond the invisible wall can be received at the base station if collisions in clusterhead candidacy messages, allowing motes within the invisible wall to become clusterheads. Motes that did not get that message could then elect to be a clusterhead, allowing messages in from beyond the wall. Starvation does not happen in LEACH since the clusterhead elections happen independent of each other, as shown in Figure 3.6 which is created from the same network with the same configuration settings using LEACH as an election protocol. Therefore, a solution is needed to avoid creating the invisible wall that the Clusterhead Candidacy phase has created.

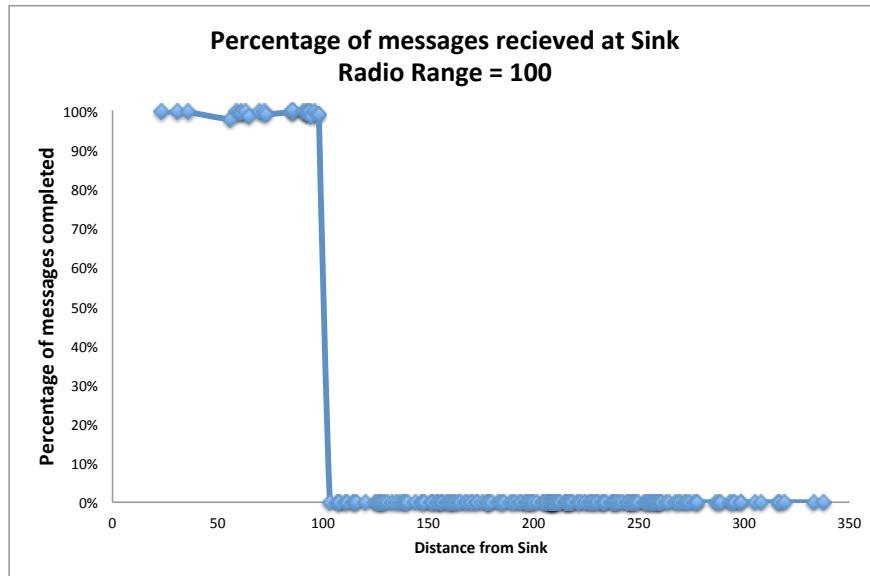


Figure 3.4: An invisible wall created due to HCCP clusterhead candidacy messages.

There are two different approaches to eliminate the problem of the ‘invisible wall’ in HCCP.

1. The sink can sleep for a cycle. Since the sink would not send a clusterhead

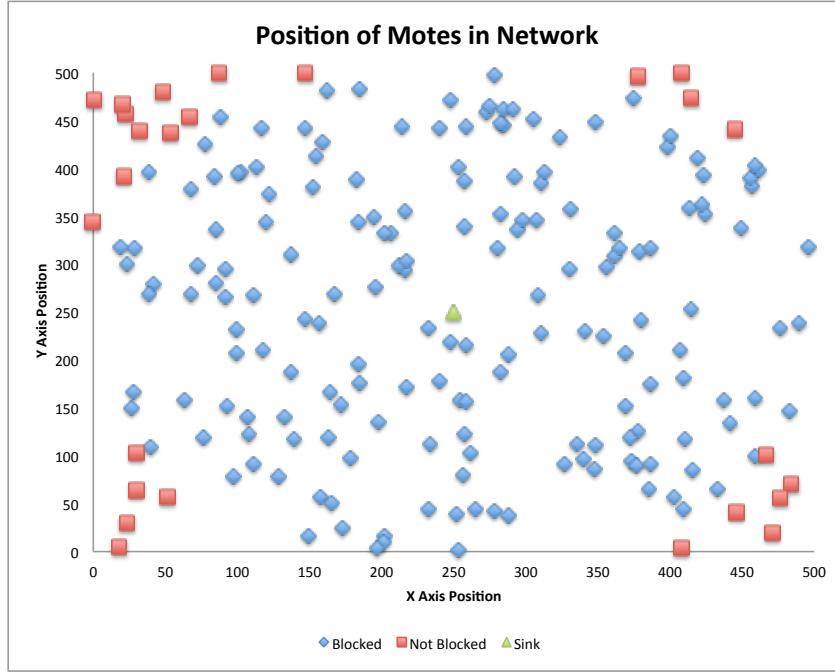


Figure 3.5: HCCP messages received at the base station. There is no way messages outside the radio range of the sink can send messages to the sink.

candidacy message, the surrounding motes would then be able to become clusterheads. This would allow messages to flow into the motes in range of the sink from motes out of range of the sink.

2. Motes can choose to be clusterheads *even if* there is a better clusterhead in range. Since the sink will always be a better clusterhead than a mote, a mote should be able to ignore all other candidacy messages and become a clusterhead.

Both of these have drawbacks. If the sink turns off for a cycle, the latency time for messages will go up, since no messages will be received at the sink for this time. If clusterheads are allowed to randomly choose to be clusterheads, there will be more collisions due to the increased number of messages being sent during clusterhead

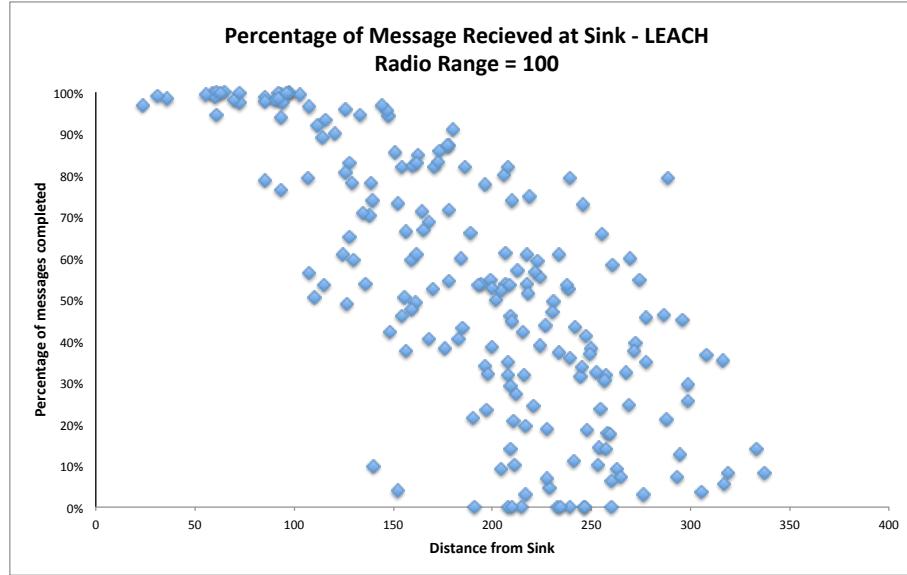


Figure 3.6: LEACH messages reached at the base station, note there is no ‘wall’ blocking messages.

messaging times (clusterhead elections and the like).

To discover which of these methods is better to use, a simulation of both LEACH and HCCP was created and the results were compared, and are compared later in section 4.1.5.

3.2.2 HCCP Goodness Delay

To ensure that motes that are better at being clusterheads are chosen to be clusterheads, the timing that is inherent in an election/run WSN such as LEACH is leveraged. If a mote has the right qualifications for being a good clusterhead, it announces its candidacy first. If another mote that was going to announce its clusterhead candidacy hears the a different mote’s candidacy message first, it will not be a clusterhead that round.

The Goodness Delay is created by delaying a percentage of the available time.

Motes that are well suited for being a clusterhead will create a small delay (therefore announcing itself as a potential clusterhead earlier), and motes that are not well suited for being a clusterhead will have longer delays.

Since the heterogeneous assessment can focus on different heterogeneous factors, the way of determining how long to delay must be robust to the possible focuses. To reflect this, a percentage of time to delay is put into a weighted average of all the factors that will contribute to the delay timing. The code looks like the following:

```

private double getGoodnessDelay()
{
    double delay = 0;

    // add up to n% of available time based on battery
    delay = (availableTime * BATTERY_POWER_WEIGHT *
              (1 - battery.getPercentLeft()) );

    // n% of available time based on mission
    // delay gets longer if this node wants to be a clustermote
    delay = delay + (availableTime * SENSOR_MISSION_WEIGHT *
                      sensorMission );

    // n% on messagequeue size
    delay = delay + (availableTime * MESSAGE_QUEUE_WEIGHT *
                      ( messageQueue.size()/(double)maxQueueSize ) );

    // n% of random
    delay = delay + (availableTime * RANDOM_WEIGHT *
                      (goodnessRM.nextDouble()));

    // n% of duty cycling based on last round as CH.
    // more delay if I was just a clusterhead.
    delay = delay + (availableTime * DUTY_CYCLE_WEIGHT *
                      ( 1- Math.min(1, lastRoundAsClusterhead * chanceOfBeingCH)) );

    return delay;
}

```

Continuous Nature of HCCP Goodness Delay

Since HCCP uses a Goodness Delay to advertise the heterogeneity of the network, it can be considered a continuous statistic. This is as opposed to the two/three/n-tier network that describes a network as a binary relationship between motes: a mote is a routing mote or not. HCCP takes what is known about the mote, and translates that information into a continuous distribution via the Goodness Delay algorithm. The Goodness Delay calculates a percentage of ‘Goodness as clusterhead’, and maps that Goodness to an amount of time to delay before becoming a clusterhead. Therefore, HCCP considers every mote to be a heterogeneous mote, even in a homogeneous network. Even in a homogeneous network, not all motes will be the same. Some motes will have batteries that are drained more than others, and some will have no space in their message queue to become a clusterhead. This approach considers such factors.

Drawbacks to Goodness Delay

When a network is first deployed, all motes will have full batteries, empty message queues and will generate approximately the same Goodness Delay. This would cause the Clusterhead Candidacy and Clusterhead Election time to effectively only be in the first 10th or so of the time allotted to the stage.

Conversely, if motes are added to the network after it has run for a while, the HCCP Goodness Delay will use much more of the stage time to announce clusterheads, since the new motes will have new batteries and empty message queues and therefore have short Goodness Delays, while the old motes will have longer Goodness Delays.

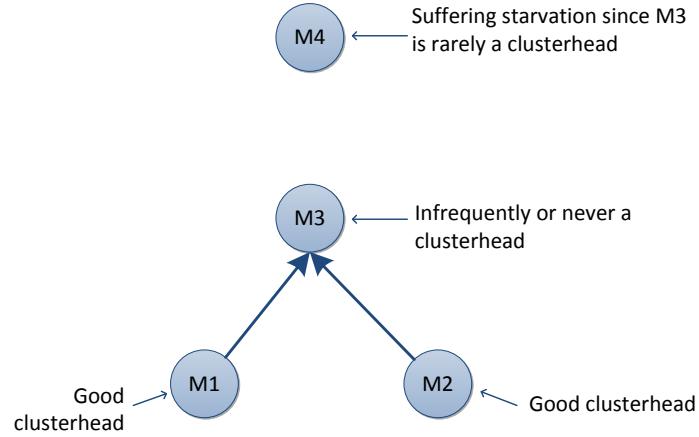


Figure 3.7: M3 never becomes a clusterhead since M1 and M2 announce their Clusterhead Candidacy first. Since M3 is never a clusterhead, M4 never has a chance to send its messages through the network.

Due to this, HCCP is very robust to changes after the network has been run for any amount of time.

Suboptimal Clusterheads

Suboptimal clusterheads are crucial to the success of a network. Starvation can occur in any place in the network if a mote is surrounded by other motes that are also poor clusterheads, as shown in Figure 3.7. Surrounding motes are never clusterheads since the surrounding nodes are in radio range of a good clusterhead. The good clusterhead will advertise itself as a clusterhead, which stops all surrounding motes from becoming clusterheads. Since all the surrounding motes are always clustermates, messages never leave the mote suffering from starvation. Since this starvation can

take place anywhere in the network, it could be called in-network starvation.

To avoid in-network starvation, all motes should occasionally become clusterheads, even if the mote is not well-suited to becoming a clusterhead. Since the mote is not well-suited to being a clusterhead, this is called becoming a suboptimal clusterhead. Suboptimal clusterheads prevent in-network starvation by becoming a clusterhead to a mote that is suffering in-network starvation; collecting its messages and hopping the messages through the network.

Due to the timed nature of HCCP elections, suboptimal clusterheads will announce their candidacy later than good clusterheads. Since the announcement is later, only motes that are suffering starvation or have no better options (better being defined as a clusterhead that announces earlier or has a lower beacon rank for routing) will choose to follow the suboptimal clusterhead. This means that cluster sizes for suboptimal clusterheads should be smaller than a normal cluster size. If HCCP did not use a timed election, then the cluster for the suboptimal clusterhead's cluster may get too large, sending too many messages and overrun the clusterhead's message queue (too many messages received from the cluster), therefore wasting more energy than needed to be a clusterhead.

Suboptimal First-order Clusterheads

Since the sink can be considered to be a clusterhead, and is a clusterhead for each round, the motes in radio range of the sink (called first-order clusterheads) should be given a different rate of choosing to be a suboptimal clusterhead. A first-order clusterhead is a clusterhead that has a beacon rank of 1, that is, is within radio range

of the sink.

Since all messages need to be hopped through the motes that are next to the sink, they need to spend more time as clusterheads. This is so the first-order clusterheads can receive messages from the rest of the network and relay them to the sink.

Sink Sleep

Another solution to getting messages to the first-order motes, is to turn the sink off for a round. Turning the sink off will allow the first-order motes to become clusterheads. Allowing the first-order motes to be clusterheads will allow the messages from the network to flow into the first-order motes. Once the messages have flowed into the first-order motes, the messages can easily be hopped to the sink.

3.2.3 Roundtable Discussion Functions

As networks were simulated, it was obvious that beacon routing is not efficient at sharing the beacon with nodes at the edges of the network. This is due to the fact that the beacon is shared once per round by the clusterhead. In an optimal case, if motes have range 100 arbitrary distance units (where range is defined as a radius), a mote 300 units away would require 3 rounds at minimum to receive a beacon. This is illustrated in Figure 3.8. The latency is due to the hopping nature of wireless sensor networks. As a beacon is sent out, only the first nodes in range receive a beacon. After that, the beacon must be sent out from those nodes, and so on.

To reduce this latency, HCCP uses the Roundtable Discussion to disseminate this important information. The roundtable time can be used by motes to query for

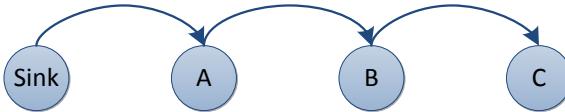


Figure 3.8: A mote range 300 away would take at best 3 rounds to receive a beacon given a range of 100.

beacons, share beacons and for clusterheads to opt out of the clusterhead role.

A solution to beacon routing is to flood the network with routing information at every roundtable. The problem with this would be the number of collisions it would create. For instance, in a dense network, if every node were to broadcast its beacon after receiving an update to its beacon, nearly all messages would turn to gibberish. To prevent this, a balance of density to number of nodes beaconing must be found, or sufficient time must be given to the motes for each to beacon without having excessive collisions, making all collided messages overhead.

There are benefits and drawbacks to either method of sharing routing information. Sharing the beacon you received immediately will cause collisions, but using CSMA can mitigate many of these problems. Checking to see if the line is clear, doing a CSMA backoff if the line is busy is a simple and effective (though time consuming) way of dealing with most of the collisions. Collisions will still occur, but the worst case is that some motes might have a beacon rank that is too high. The mote in the center does not receive a beacon from the closer motes, since the two messages collided. Two cases can occur from here, either the center mote receives a beacon from a mote further away from the sink, or re-requests a beacon from the surrounding motes and receives the proper rank.

Overall, the Roundtable Discussion can be a simple exchange of CSMA messages,

or can be made to do specialized tasks as needed by the network. The simulations and testbed deployments will share routing information using simple CSMA messages.

Options for Tuning the Roundtable Discussion

The Roundtable Discussion time is very expensive on the network, as all the motes are on during this time. If a long time is chosen for the Roundtable Discussion, the network life will be negatively effected. If a long-lived network is desired, then the Roundtable Discussion should be left out all together, as long-lived networks are purpose built for lifespan and do not need the energy draw of the Roundtable Discussion.

If no extremely important information needs to be sent during the roundtable time, only a subset of the motes need to be online to flood the message across the network. This too, will increase the lifespan of the network.

3.2.4 Clusterhead Choice Timeout

Consider the situation illustrated in Figure 3.9 where a clustermote B has a lower beacon rank than a nearby ‘very good’ clusterhead. The mote should choose the good clusterhead as its clusterhead over the other potential clusterheads that may have a better beacon rank, but are significantly poorer clusterheads.

The question that must be answered is ‘what is significantly poorer?’’. A way to solve this problem is to use the ‘Goodness Delay’ timing that is already inherent in HCCP. Once a clustermote has heard a clusterhead announcement, there is a timeout clock that starts. If another clusterhead makes an announcement before the timeout

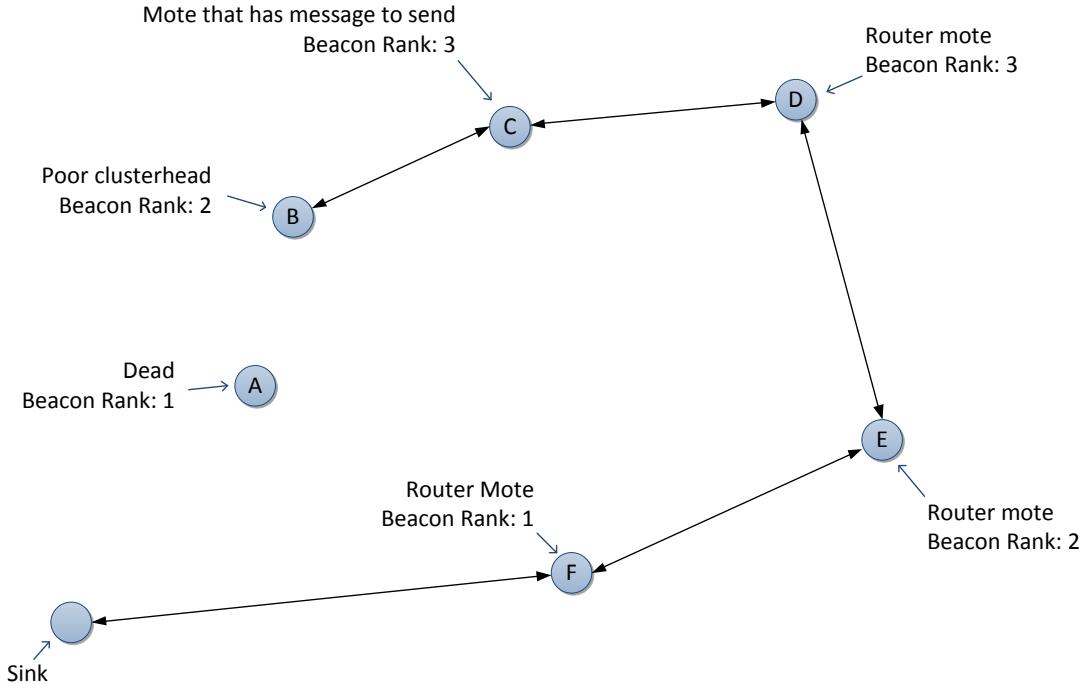


Figure 3.9: Mote C would choose mote D as a clusterhead even though mote B is closer to the sink.

expires, and that new clusterhead has a lower (better) beacon rank, the mote should join that new cluster. Once the timeout has expired the quality clusterheads that are announcing are considered significantly poorer than a known clusterhead. This is due to the ‘Goodness Delay’ that HCCP uses to discover heterogeneity in the network. So, despite a mote having a lower beacon, it might not be the best clusterhead, and the clusterhead that is good should be used.

Chapter 4

Experimental Setup and Results

Simulations and a testbed deployment were created to show the benefits of using HCCP over LEACH and were used to compare the two protocols. To compare the performance of HCCP and LEACH, simulations were created for both of them using custom simulation software. The simulation software is freely available on Github at (<http://github.com/robguderian/hccp>). The simulation provides a high-level view of a WSN, focusing on how the network can function together and give insight into the network, allowing many different factors to be logged that would otherwise be impossible to track.

For the testbed deployment, both HCCP and LEACH code were created for four different types of WSN motes that have the ability to communicate with each other. A small deployment was then run to help show that the simulations are accurate and that HCCP can work in a real-world environment.

4.1 Simulation of HCCP and LEACH

A custom simulator was developed to simulate the running of HCCP. Other network simulation suites are available, such as OmNet++ [32] with its related suites such as Castalia [33] are widely used, and provide tools for analysis. These tools were tested, and were deemed unfit for the desired simulation setups. Tracking the number of messages sent, from where, route taken and number of times a given message has been received at the sink are possible, but difficult to collect. OmNet++ messages can only be in one mote's message queue at a time since it can only have one 'owner'. Since a message can only have one owner, it is very difficult to track how many times a given message has been received at the sink. Also, the MAC layer protocol is not simple to change while running a simulation. LEACH uses both CSMA and TDMA MAC protocols, so to properly simulate LEACH either switching the MAC protocol must occur, or the finer details of the protocols must be abstracted and be viewed simply as access to a radio.

WSN code can be written in such a way that the physical layer can be largely ignored. The physical layer controls the radio, modulating the frequency to communicate with surrounding motes. This includes what frequency or protocol (such as 802.15.4 or ANT radios) the motes use. The simulator abstracts the physical layer, as it is not the focus of the simulation or research. It is assumed that the radios work, have a given range and draw power when on.

When building the simulator, the problem of simulating a WSN was viewed as a queuing theory problem more than a networking problem. In doing this, many of the network problems are abstracted away, such as radio channels or how collisions can

be recovered from. HCCP currently only uses one radio channel, so only one channel was created. Messages that have collided are assumed lost, and unrecoverable.

With these assumptions and abstractions, a simple to understand simulation could be made to collect important statistics about the inner workings of the network.

Simulator Capabilities

The custom simulator is able to create networks of any size with a two dimensional rectangular space in which to place the motes. Attention was given to the use of random number generators and randomly created events. Random seeds are kept, and random events can happen at the same times in networks that are being compared.

Each mote is given its own random number generator with its own seed to generate random events. Separate random number generators are used to draw random numbers for all tasks, ensuring the simulated motes generate the same random numbers for both the LEACH and HCCP simulations.

The package used to generate random numbers, and to maintain the event queue was SSJ [34]. SSJ is a well-respected Java simulation package. The pseudorandom number generator and the event queue were built for simulating queuing theory events, and therefore were easily applied to simulating events in wireless sensor networks.

SSJ provides excellent facilities for collecting statistics on the motes. SSJ can collect continuous data, such as calculating what percentage of the time a given mote was on or off; or, collect discrete events, such as number of messages received, sent and lost.

Design of Motes in the Simulation

The custom simulator simulates motes that are an abstraction of motes in real life. Motes have been abstracted to have the following features:

1. **A battery:** The battery drains faster when the mote is on, and slowly when the mote is sleeping. Motes can draw more power if they are power intensive motes, or less power if they are power efficient motes. The batteries can start with more or less energy, which simulates having a larger or smaller battery. Since batteries in reality do not always have the same amount of battery power, the simulator applies jitter to the amount of power given to the motes. This amount is configurable, and generally assumed to be quite small.
2. **A message queue:** The queues contain messages created by this mote, and messages that have hopped into this mote. The queue has finite space. If the queue is full, the message that should be added to the queue is lost.
3. **Sensor readings:** Sensor readings happen at a specified frequency. Sensor readings get turned into messages which are added to the message queue. If the message queue is full, the message with the reading is lost.
4. **A simple MAC layer:** The MAC layer is designed to be simple and abstract. CSMA uses backoffs of the channel is currently being used. CSMA waits a random amount of time before attempting to send the message again. Each mote needs its own random number generator to generate a random backoff time.

5. **A radio:** All devices have radios that have a given range. If a neighbouring mote is within range of a device, the two motes can communicate. All radio links are assumed to be one-way links, as some radios could have more transmission power than other radios. The units of the range are arbitrary distance units, and could be interpreted as centimetres, metres or even miles.
6. **A position:** Motes are given a position on the two dimensional plane. Any mote can be set to be mobile in the network, and can move at any time. This movement will change which surrounding motes the moving mote can communicate with.

This is consistent with the high-level description Akyildiz et al. [1] gave to describe a mote, and with Figure 1.1.

Motes can be of various different types: basestations, routing nodes, or sensor nodes. Basestations are the sinks, once a message is received at the basestation it is marked as completed. Routing nodes are motes that have no sensors, and are therefore ideal for routing messages. Sensor nodes are motes that make sensor readings, and the sensor readings are turned into messages. The type of sensors and number of sensors have been abstracted away. Basestations are assumed to have wall power, sensor and routing nodes are generally assumed to have battery power, but the simulator has the capabilities of giving them wall power.

Heterogeneity is given to the motes by initializing the motes with different properties, such as:

- more or less battery power

- more or less power drawn when on or off
- more or less available space for a message queue
- longer or shorter radio transmission range
- the ability to move or not and how fast the movement is
- frequency of sensing and creating messages

The number of sensors or type of sensors can be abstracted to being a frequency of sensing events. If a mote has many sensors, it can be abstracted to have more sensing events. A routing mote is the other extreme, in that a routing mote has no sensing events and therefore no frequency of sensing events.

All motes in the custom simulation use common code for the MAC layer, as this is a requirement for heterogeneous WSN. Motes in heterogeneous WSNs have differing hardware, energy levels and capabilities, but share a common mode of communication.

Motes in the Simulations A set of motes were reused for consistency across all the simulations. This provides standardization of the motes across the simulations while providing heterogeneity across the simulations. The power draw of all the motes were kept the same across all the types of motes for all simulations. The motes varied in queue size and battery size.

- **Router** - A special type of mote that does not have any sensors. It is given a large queue and a large battery making it ideal for routing messages.
- **Normal** - An average mote. Most of the motes in the network are set to this kind of mote. It has an average battery, and an average queue.

- **Expensive** - The expensive mote simulates a mote that has expensive sensors that are power hungry devices. This mote has a larger message queue than the normal mote, but a smaller battery.
- **Very Expensive** - The very expensive mote simulates a mote that has more power hungry sensors than the Expensive mote, a smaller battery, and a smaller message queue.
- **Super Expensive** - The Super Expensive mote simulates a mote that has the most expensive sensors that put a very large load on the battery. It therefore has the smallest battery, and the smallest message queue.

4.1.1 Modifications to LEACH for Simulation

LEACH does not specify which routing algorithms should be used when using LEACH. For comparison reasons, two routing algorithms were used in creating LEACH: beacon routing and preset routing, where every node knows its beacon rank from the beginning of the simulation.

In the beacon routing implementation, the mote published its beacon rank when it announced it was a clusterhead, during the clusterhead announcement phase. This is not very efficient, but is a realistic way of disseminating the routing information to the network.

Since HCCP has facilities to efficiently disseminate routing information, a preset routing simulation was created. The preset routing simulation sets each of the motes appropriate beacon rank before the simulation starts. This permits the network to

work as if the network has been alive for long enough that the routing information has been disseminated throughout the network.

4.1.2 Modifications to HCCP for Simulation

HCCP has facilities for routing, but does not specify which routing algorithm should be used. Because a routing algorithm is needed for the network to work once it is deployed, the simulation has been implemented with beacon routing with HCCP, as HCCP will then be evenly comparable to LEACH. The beacon information is disseminated during the roundtable discussion, as prescribed by HCCP.

Since LEACH has been provided preset routing, the HCCP simulation has also been given preset routing, so it can be fairly compared to LEACH. Simulations with the preset routing tables were done, as were simulations without preset routing tables to show the differences in data dissemination between the networks.

4.1.3 Observations and Results

The worst-case baseline runtime for a mote in the simulation setup is 40,776s (the mote is always on), and the best-case runtime is 256,200s (the mote is only on during the mandatory times). Clearly the best-case is impossible to reach, as all motes would be sleeping for the entire cluster runtime (TDMA schedule runtime), never sending any data messages. Since the motes are asleep for the entire cluster runtime, it will also will never send hop messages, which means the mote will never have its messages reach the sink and is therefore useless. If a mote is always a clusterhead, it will be on for the entire simulation making the worst-case runtime a possibility in the

simulation. Duty-cycling ensures that a mote will never only run as a clusterhead.

All permutations of the possible Goodness Delay factor configurations were run to provide insight to which HCCP parameters that build the Goodness Delay (described in section 3.2.2). The permutations were run as the following: 100% Battery power; 99% Battery power, 1% Message queue; 99% Battery power, 1% Random; and so on. Each parameter contributes a percentage of the time that the mote waits before it announces its Clusterhead Candidacy or Clusterhead Announcement. Since there are 5 factors, each are considered separately. The LEACH-style clusterhead percentage is controlled separately from the 5 HCCP factors and are also viewed separately. Motes will elect to be a Clusterhead Candidate using the same equation as LEACH uses for electing clusterheads.

The x -axis of the following graphs show the percentage that the listed HCCP factor contributes to the Goodness Delay. The remainder of the Goodness Delay is a permutation of the remaining HCCP factors. For instance: if the contribution to the Goodness Delay of Battery Power is 50%, the other 50% will be made of up all the possible permutations of Random, Duty Cycle, Message Queue and Sensor Mission. The charts all have the same x and y axis to make comparisons clearer.

HCCP, much like LEACH, can be set up to vary the percentage of clusterheads per round in the network. Since motes that are clusterheads use more power per round than motes that are not clusterheads, choosing a higher percentage of clusterheads per round causes the network to have a shorter lifespan. This effect can be seen in Figure 4.1, as the percentage of clusterheads increases the lifespan of individual motes drops. There is a peak at 3% clusterheads for average lifespan since motes that choose

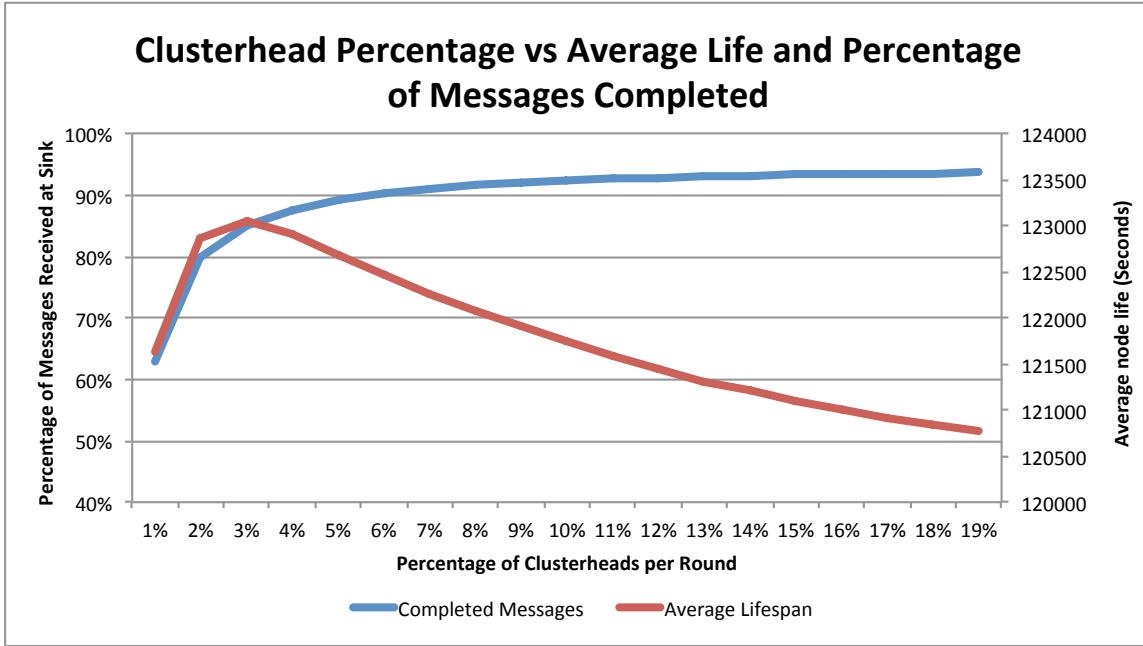


Figure 4.1: Visualization showing the relationship between Percentage of Clusterheads per round and simulation results.

clusters turn off for the remainder of the HCCP phase. These small energy savings add up to make the average lifespan of the network longer. Heinzelman et al.[21] also saw the same phenomenon happen while testing LEACH, finding that the motes dissipate the least energy in their test case at about 5% clusterheads per round. Since HCCP uses many of the design elements of LEACH it is not surprising that the two have similar optimal percentages of clusterheads per round.

The message throughput increases with the percentage of motes that are clusterheads each round. This makes sense, as the non-clusterhead motes would have a good selection of clusterheads to choose from, and could choose a clusterhead that is nearby to decrease the chances of a message collision. However, as the percentage of clusterheads increases, the lifespan of the network decreases. This causes a tradeoff

for a network administrator to choose from, as some networks might require longer lifespans, while other networks may value message throughput.

Effect of Goodness Delay on HCCP

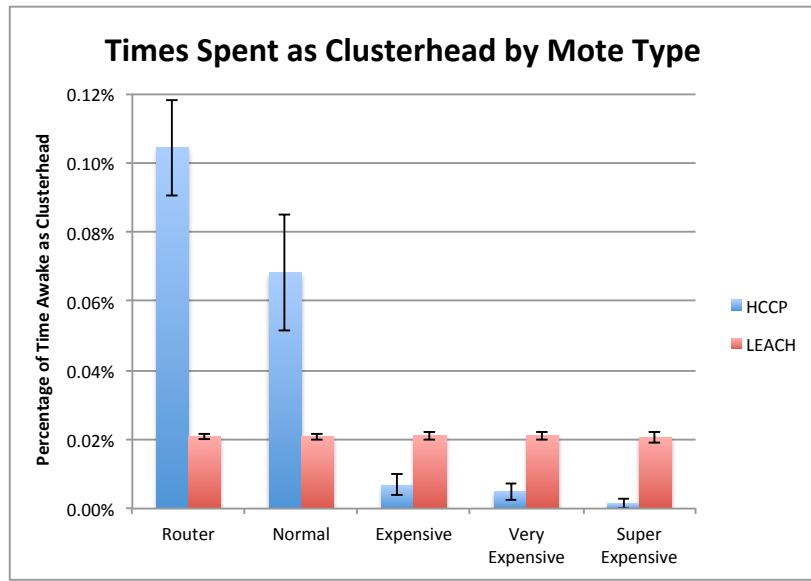


Figure 4.2: The Goodness Delay focuses the clusterhead task on motes that are suited for the task.

The Goodness Delay in HCCP is intended to make motes that are well-suited for the task clusterheads more often. Figure 4.2 shows that the Goodness Delay is effective in doing this. As the motes get more expensive, hypothetically with more power-hungry sensors, the mote gets chosen to be a clusterhead less frequently. This means that the mote will draw less power, since it is not being a clusterhead as frequently.

Note that in Figure 4.2 the percentages are quite low, as the Y axis is the percentage of time the mote was a clusterhead over the entire time it was alive, which includes long sleep periods between cycles of the protocol.

4.1.4 Discovering how Heterogeneous Factors Effect the Network

Simulations were run with all the possible permutations of the factors available in the simulation. The results show which heterogeneous factors are valuable for improving the WSN. The factors that could be used in clusterhead elections in HCCP are: residual battery power, available message queue size, sensor mission, when the last time this node was a clusterhead was, and a random variable. Breaking the problem down to the separate factors, the value of the factors can be compared.

All the permutations of the HCCP factors for Goodness Delay were run on the same network with the same random seeds. This set up the network with the same starting and running parameters so that events would happen at the same times in the various networks, allowing the networks be comparable with the given parameters.

Effect of Available Queue Size on HCCP Goodness Delay

Focusing on available queue size for a method of describing how good a mote would do as a clusterhead is an obvious choice, as clusterheads will be collecting messages from the surrounding motes and need the ability to store all the messages. If the message queue is full on a mote, it could not store the new incoming messages, therefore losing all incoming messages. If incoming messages are lost it not only is negative to the network in terms of information loss, but also wastes the energy of the poorly chosen clusterhead and all the motes that register with that cluster.

The effects of focusing on Message Queue size can be seen in Figure 4.3. As the focus on message queue increases, the average lifespan of the network drops. This is

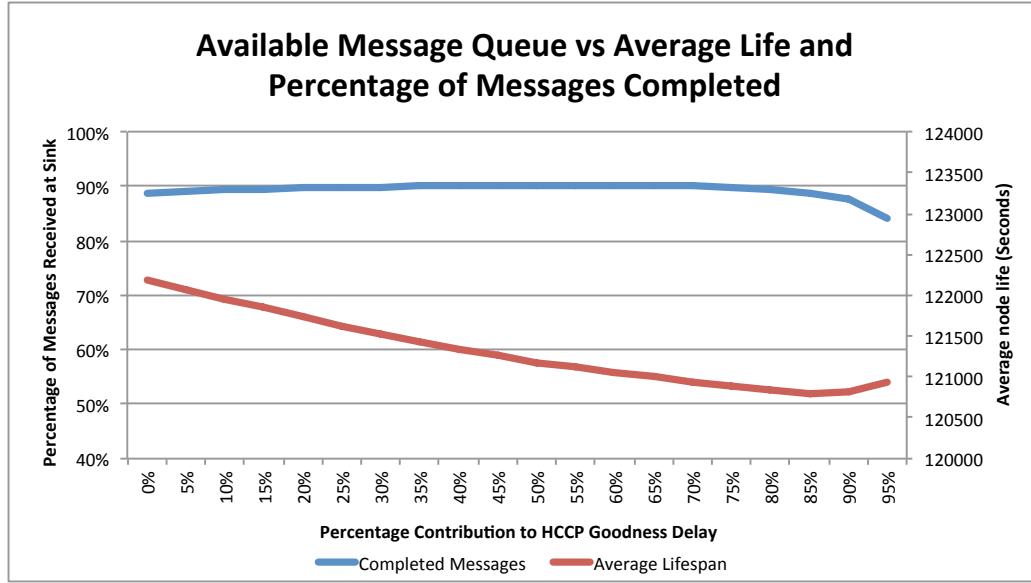


Figure 4.3: The relationship between HCCP Goodness Delay using Available Queue size and simulation results.

not surprising as the focus on battery power and duty-cycling reduces, making the network solely focus on choosing clusterheads that have large available message queue size regardless of available power. Because of this, Message Queue size is likely a good secondary factor that would control less of the Goodness Delay, balancing a factor such as Battery Power.

As the network gets busier, the chances of any mote having a full message queue is greater. When message queues start filling up, the importance of focusing on message queue size as a HCCP Goodness Delay factor increases. Because HCCP can focus on which motes can be effective clusterheads, better clusterheads will be elected. The results of this can be seen in Figures 4.4 and 4.5, which have the same x and y axes to show the difference between the two results. HCCP can handle a network that is overloaded with messages quite well, because motes with more available space to

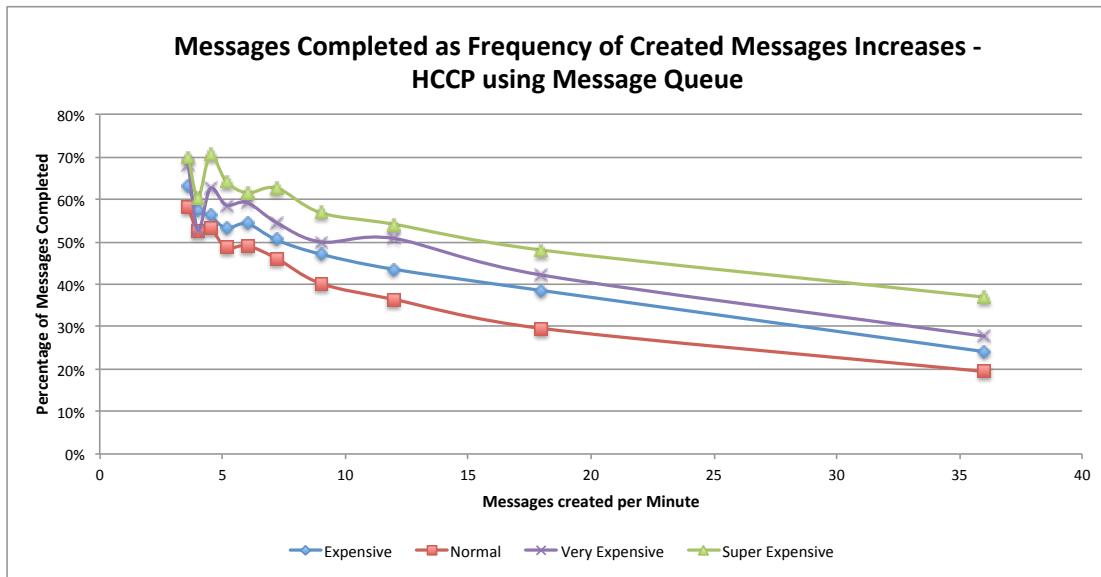


Figure 4.4: HCCP handles being overloaded with messages well, as motes with larger available message queues will opt to be clusterheads.

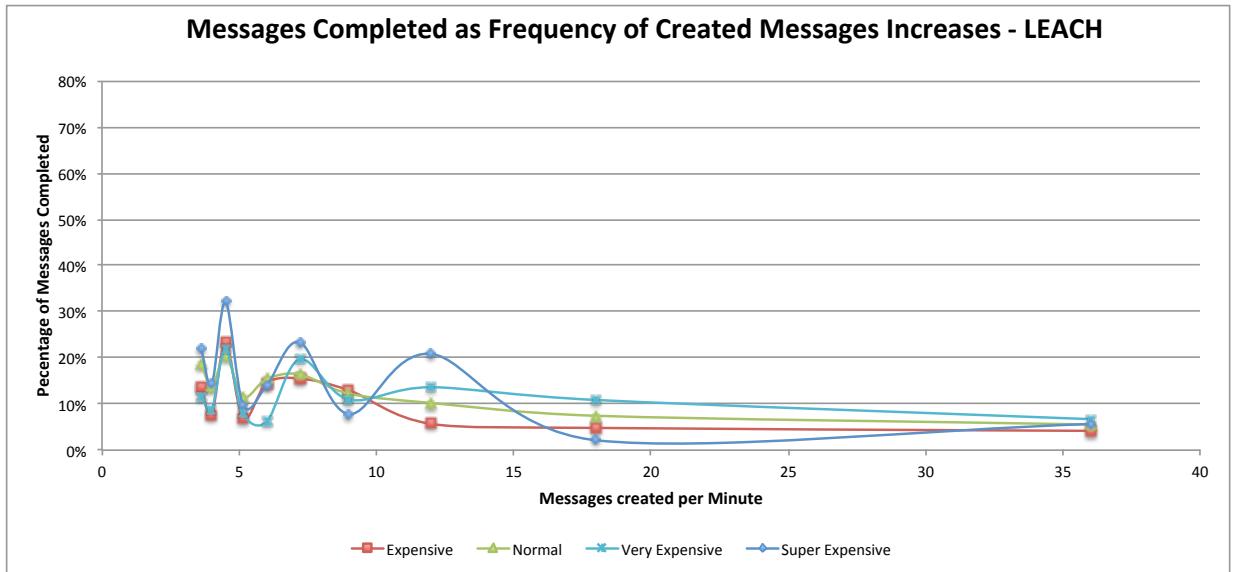


Figure 4.5: LEACH does not handle being overloaded, as it does not consider message queue size when electing clusterheads. Note the y axis is the same as Figure 4.4.

hold incoming messages from the clustermotes will be more likely to be clusterheads.

The standard deviation of the simulation is approximately $\pm 20\%$ for both HCCP and LEACH for each frequency level. This shows that HCCP focusing on Message Queue has a significant gain over LEACH in terms of message throughput for Super Expensive motes even at the most overloaded network run. It is also interesting to note that the signal-to-noise ratio for LEACH is approximately 100%, even at the least overloaded simulation run.

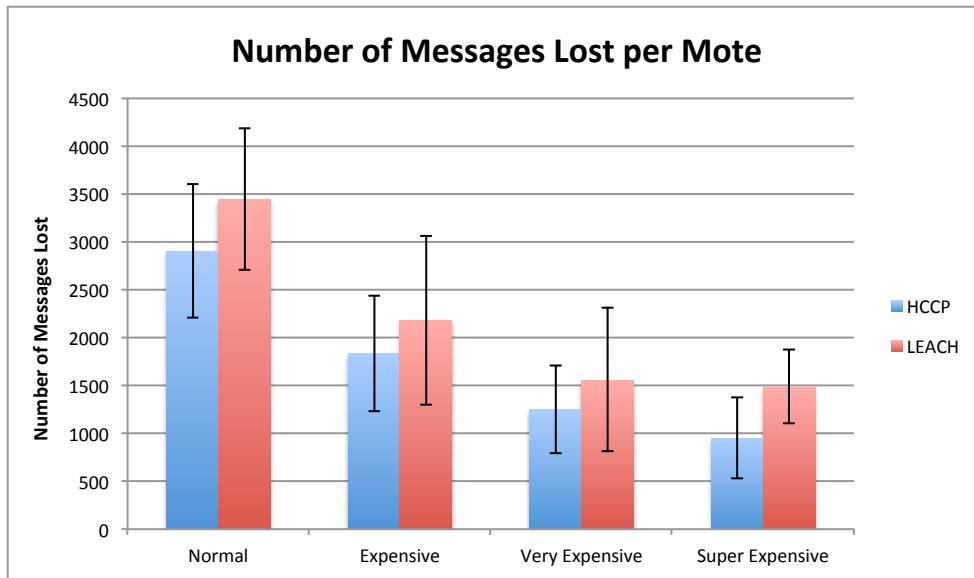


Figure 4.6: HCCP loses fewer messages due to full message queues than LEACH.

HCCP focusing on Message Queue will have the effect that fewer messages will be lost due to having no available space in the message queue. Fewer messages will be lost as motes that have more available message queue space will opt to be clusterheads, which will then free up space in the surrounding motes to use. Figure 4.6 shows that HCCP does have fewer lost messages, significantly more for Super Expensive motes. This shows that HCCP has gains over LEACH in how many messages are lost due to

full message queues.

Effect of Available Battery Power on HCCP Goodness Delay

Choosing motes with more available battery power creates network with long average lifespans, as seen in Figure 4.7. Comparing Figure 4.7 to Figures 4.1 4.3 4.10 and 4.11 it is clear that focusing the HCCP Goodness Delay on available battery power is second only to Sensor Mission without the drawback of having to know any information about the network before deployment.

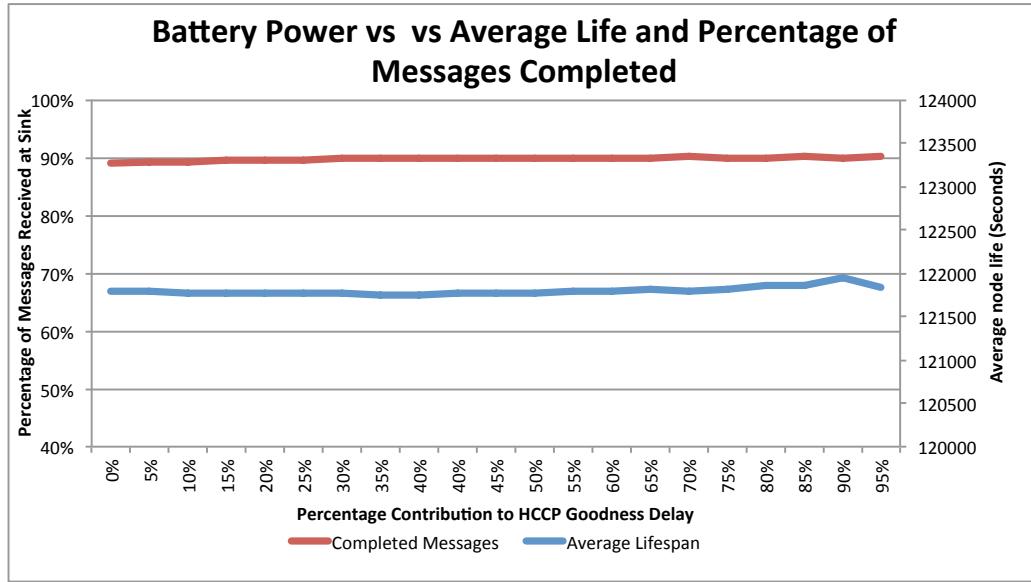


Figure 4.7: The relationship between HCCP Goodness Delay using Available Battery Power size and simulation results.

An interesting result of focusing on Battery Power is that fewer messages are lost due to motes dying. This makes sense, since as the battery in the mote nears depletion, the mote will not elect to be a clusterhead. Because the mote is not a clusterhead, the mote is not accepting messages from surrounding motes, rather it is

getting rid of all the messages in its message queue. Then, once the mote finally dies, its message queue will be as empty as possible and minimal messages are lost.

This effect can be seen in Figure 4.8. The election protocol in LEACH is naive, electing clusterheads at random, causing motes that are nearing death to become clusterheads. Notice that there is a high variability in how many messages are lost in dying motes in LEACH. In fact, the simulation results showed that there were motes that died with their message queues completely full. HCCP, on the other hand, has very low variability due to the network avoiding using dying motes as clusterheads.

The end result of using Battery Power as an HCCP factor, is that fewer messages are lost, therefore more messages successfully reach the sink. This can be seen in Figure 4.9.

Effect of Duty Cycling on HCCP Goodness Delay

The HCCP Duty Cycle factor causes motes to live longer, but deliver fewer messages. This is the common WSN throughput versus lifespan tradeoff. As motes focus on duty-cycling, they become clusterhead less frequently, which causes message throughput to drop. This relationship can be seen in Figure 4.10, as the focus on Duty Cycling increases, the message throughput drops while the average mote lifespan increases. This trend is followed until 85% focus on Duty Cycling, at which time the same effect as clusterhead percentage in LEACH occurs, too few motes are announcing themselves as clusterheads, causing the rest of the motes to be powered up longer waiting for clusterhead candidacy and clusterhead announcement messages. Since these motes are no longer waiting for announcement messages, the lifespan

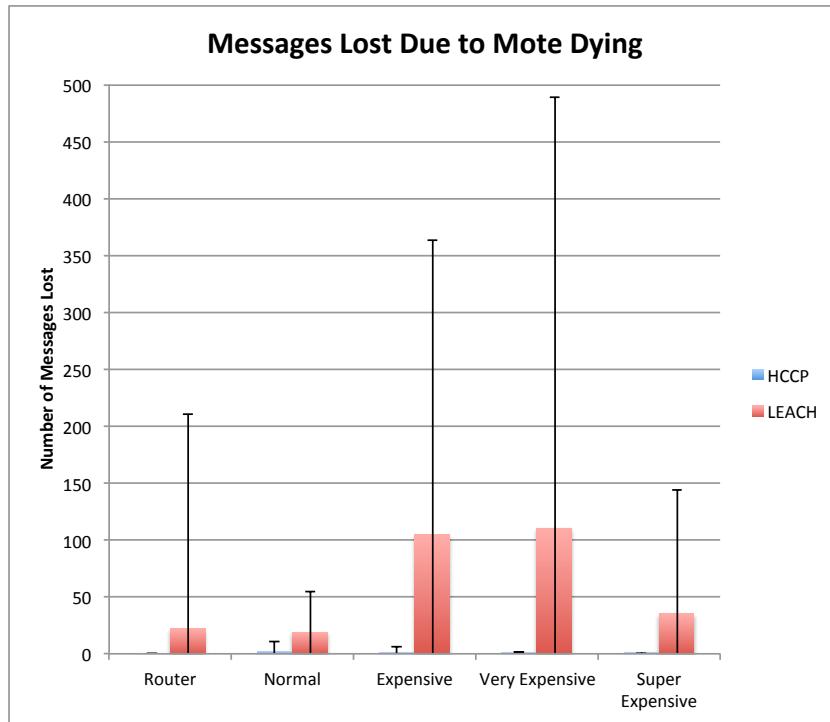


Figure 4.8: HCCP focused on Battery Power loses less messages since dying motes will not elect to be clusterheads.

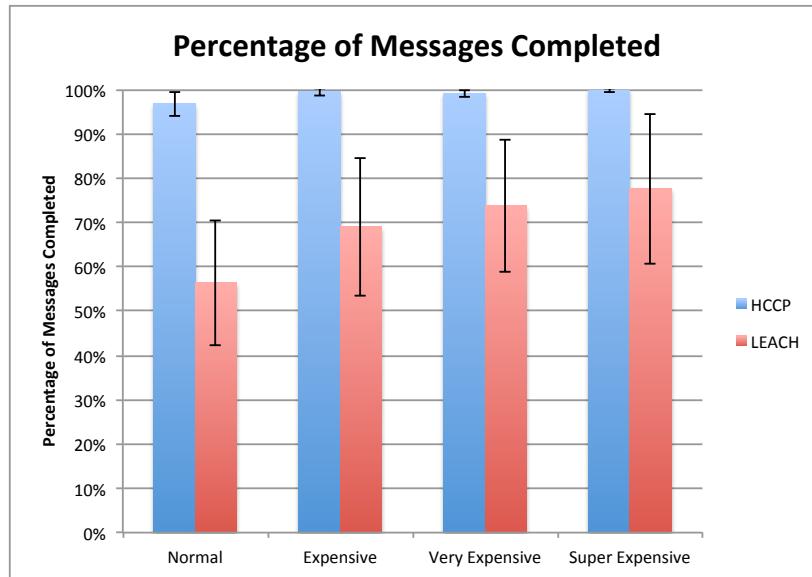


Figure 4.9: HCCP focused on Battery Power has better throughput than LEACH since fewer messages are lost in dying motes.

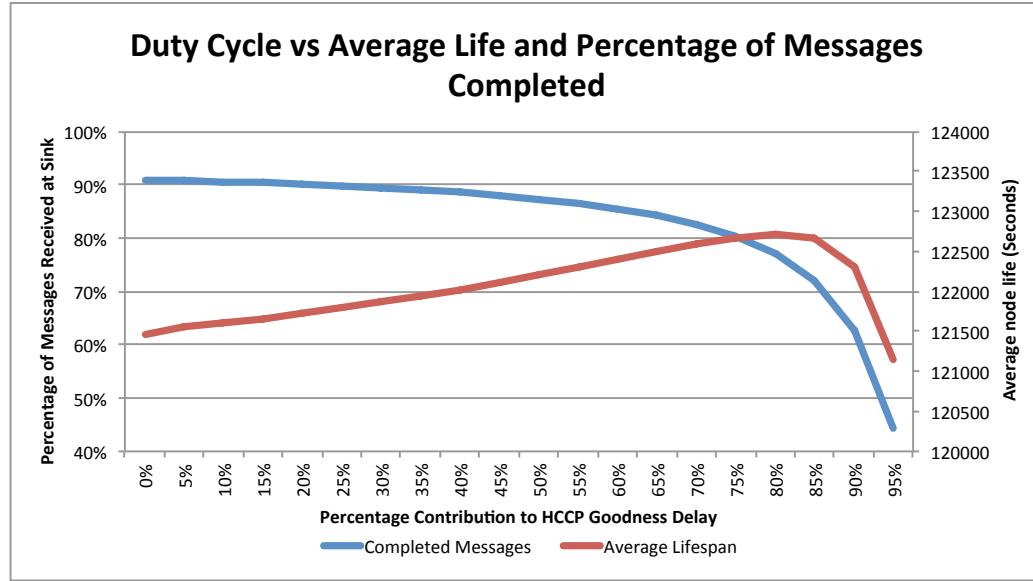


Figure 4.10: Visualization showing the relationship between HCCP Goodness Delay using Duty Cycling and simulation results.

drops. Though there are obvious gains for life span to focusing on Duty Cycling, Duty Cycling is a better secondary HCCP Goodness factor with a minority of the control of the Goodness Delay time.

The problem with Duty Cycle as the main HCCP Goodness factor is that it is not providing enough value to make up for the overhead that HCCP puts on the network. At no time did HCCP using Duty Cycling beat LEACH in simulation. LEACH has duty cycling built in to the design. HCCP uses that same duty cycling when choosing if the node should be a clusterhead, then describes how good the clusterhead is by announcing the clusterhead sooner or later. HCCP Duty cycling doesn't add this value to the network, it only delays motes longer from announcing their Clusterhead Candidacy if the mote has been a clusterhead recently. Using Duty Cycling in an HCCP Goodness Delay is therefore useless and should not be used because it costs

more to run this redundant idea than the benefit it provides to the network.

Effect of Random Delay on HCCP Goodness Delay

Randomness is the backbone of LEACH, in that each node creates a random number to decide whether or not it should be a clusterhead. This factor in HCCP's heterogeneous election is just a random number that will decide how long the mote will delay before transmitting its candidacy message or clusterhead announcement. Using only the Random HCCP Goodness factor, HCCP degrades to a slightly mutated form of LEACH, as LEACH draws a random number that decides how long the mote will delay before announcing itself as a clusterhead. HCCP takes this idea, making the moderate change that motes will delay the random amount of time until it either hears a better mote (better being defined as a mote that transmits its clusterhead candidacy earlier), or gets to transmit its candidacy message.

Using Random as the sole HCCP Goodness factor does not work well, as seen in Figure 4.11 Randomness has very little effect on the network throughput or average node life. But Random works well paired with other factors as a tie-breaker, making one mote announce it's clusterhead status before a neighbouring node of similar configuration. Using Randomness as a tie breaker works especially well early on in the network's life, as many of the motes will have empty or nearly-empty messages queues and full battery power. If two motes have the same Goodness Delay they will transmit their candidacy or clusterhead announcement at the same time, potentially creating collisions in the network.

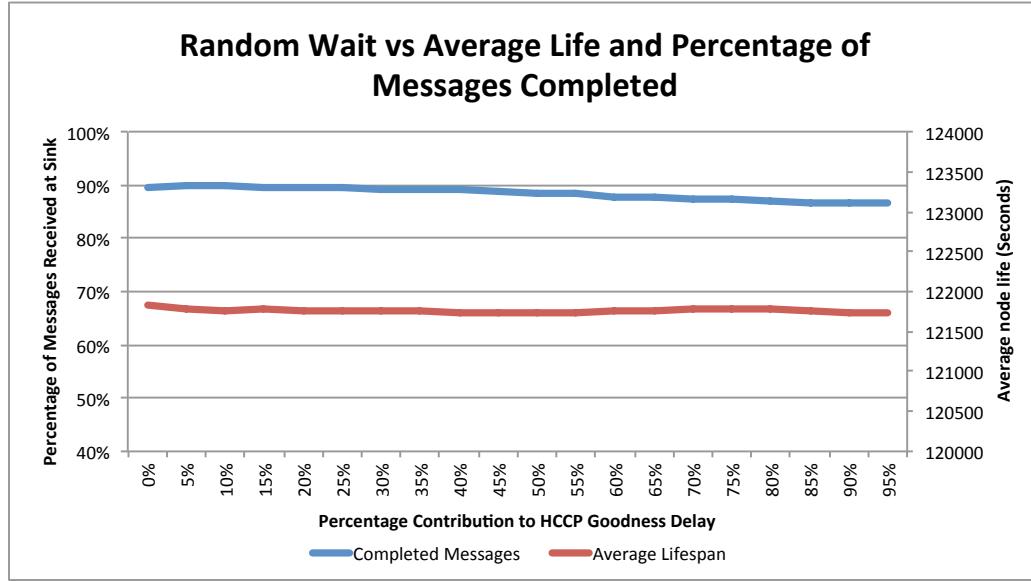


Figure 4.11: The relationship between HCCP Goodness Delay using Random Wait Times and simulation results.

Effect of Sensor Mission on HCCP Goodness Delay

Sensor mission is a percentage of how important the sensors are for a given mote. For instance, a router mote with no sensors would have a sensor mission of 0%, since it has no sensors. A mote with very expensive sensors should almost never be a clusterhead, preserving its battery power to collect more information. This value can be manually set, or automatically created by counting the number of sensors that are being read and creating a Sensor Mission percentage from that information.

Sensor mission creates a positive trend in both average node life and message throughput, as seen in Figure 4.12. This is because routers were given a sensor mission of 0 and sensor motes were given high Sensor Mission. Providing motes with a Sensor Mission value is the best case as motes will then ‘know’ how good they are at the clusterhead task, making better router motes and motes with larger message

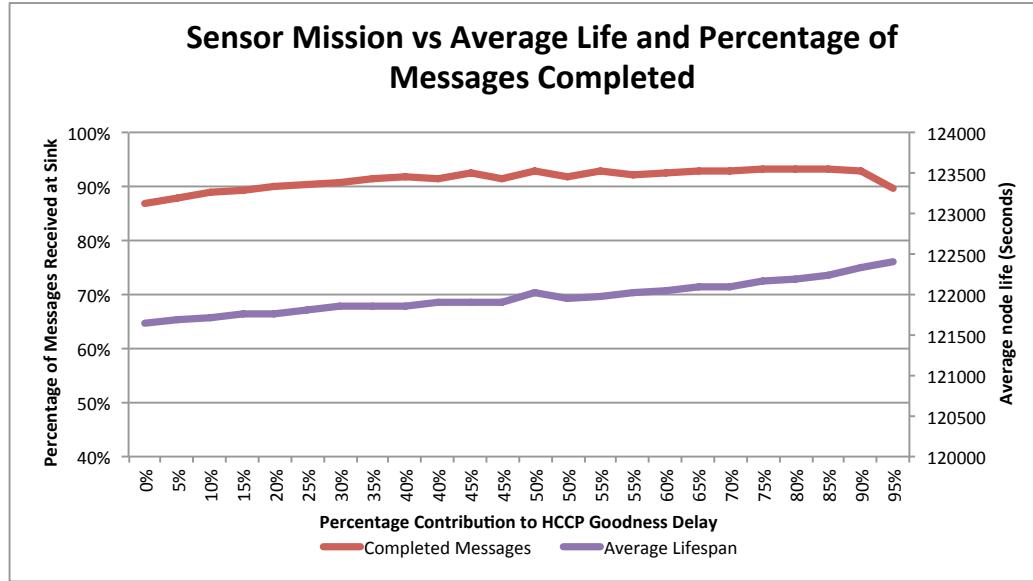


Figure 4.12: The relationship between HCCP Goodness Delay using Sensor Mission and simulation results.

queues clusterhead motes more often. The problem with providing motes with Sensor Mission values is that it is manual labour, and must be configured before the network is deployed. Also, Sensor Mission would be the most successful with a heterogeneous network; if all motes are the same and have the same Sensor Mission, the clusterhead elections would degrade to an HCCP network that is 100% focused on Random since all motes would be generating the same Goodness Delay values.

Paired Factors in Heterogeneous Elections

More than one factor can be paired together in the heterogeneous election to generate the Goodness Delay. When networks first start, all message queues will be empty and all batteries will be at 100%, this will cause collisions during the HCCP Candidate Announcements, since all motes that elect to be a clusterhead that round

will attempt to transmit their candidacy at about the same time (which will be close to immediately) if the HCCP Goodness Delay is 100% generated from a single factor. Adding a small percentage of the Random factor will vary the generated transmission times, preventing the problem of equal Goodness Delays.

When adding Randomness to the Goodness Delay calculation will first generate a goodness based on a primary factor (such as Available Message Queue Size or Available Battery Power) which makes up the majority of the time, say 90%. The Randomness will make a 10% variance of the remaining time. The variance in the delay time will effectively be a tie-breaker for the transmission times for the motes, avoiding many collisions while sending the announcements. The drawbacks to adding the Random factor to the Goodness Delay is that it adds no value to the Goodness Delay, as discussed in Section 4.1.4. Though, a small amount of Random Delay in the network is beneficial to the network due to the collisions it avoids.

4.1.5 Tuning HCCP for Power Efficiency

HCCP adds more time where all the motes are fully powered on than LEACH, because of the Roundtable Discussion and Clusterhead Candidacy stages. These extra steps cause an energy drain on the network, but add lots of value to the network in terms of message throughput and data dissemination. As mentioned in Section 3.2.3, HCCP can be tweaked to be more power efficient while still utilizing the gains from the Goodness Delay.

The easiest gain is to eliminate the Roundtable Discussion. This will improve network life, but decrease the ability to quickly disseminate information. For static

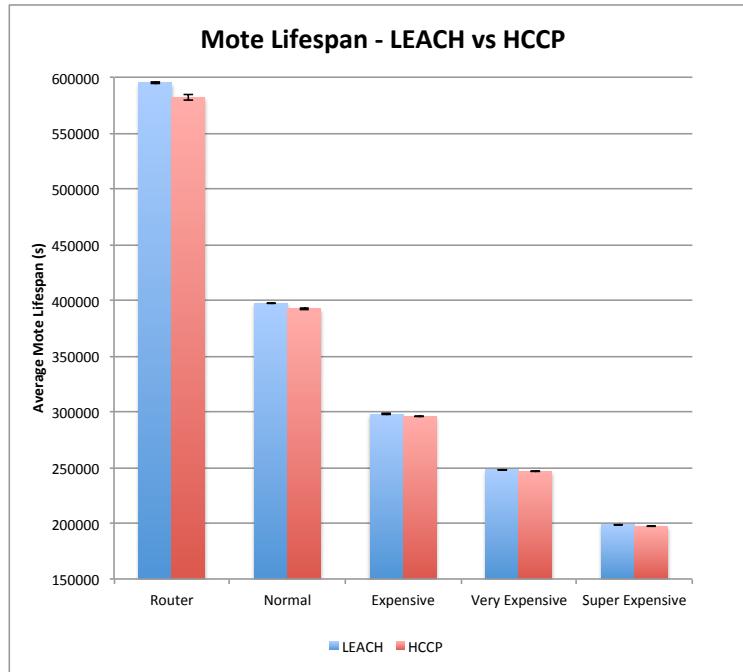


Figure 4.13: Tuning the efficiency of HCCP by minimizing Roundtable discussion allows HCCP to have networks lifespans approximately equal to LEACH.

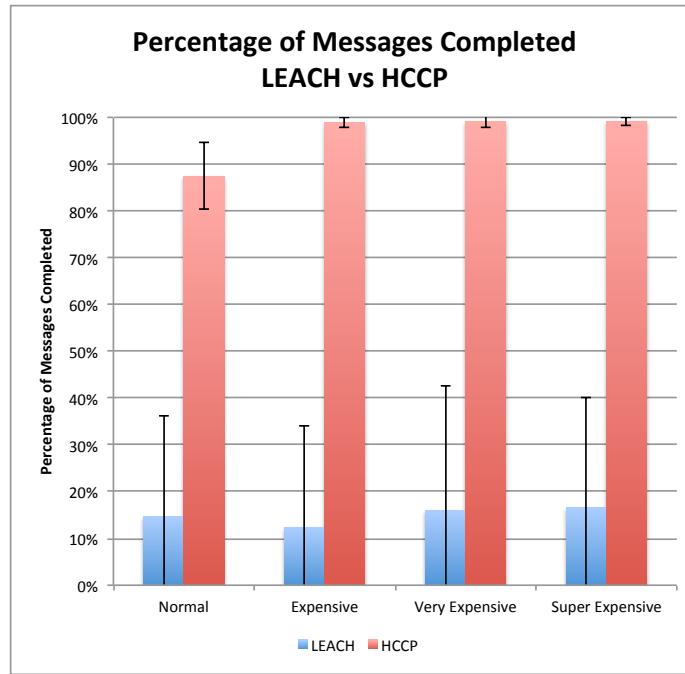


Figure 4.14: In a tuned network, HCCP still has a much higher message throughput.

networks that require a long lifespan, this is a good option.

The other option is to eliminate the Two-Stage Election, by eliminating the Clusterhead Candidacy stage. To compensate for not having the Clusterhead Candidacy stage, the Clusterhead Election stage can be used to implement the Goodness Delay features. To do this: in the Clusterhead Election, if a different mote announces itself as a clusterhead first, concede being a clusterhead and become a clustermote. This is the same idea as used in Clusterhead Candidacy, but done at the same time as a Clusterhead Election. The network will then choose slightly less optimal clusterheads, but overall creates a gain in which motes elect to be clusterheads. See Section 4.1.5 for more details about not using a Clusterhead Candidacy phase.

A well-tuned HCCP network will have a slightly shorter lifespan than a comparable LEACH network, but will have a much higher message throughput than the LEACH network due to the quality of clusterheads that are chosen. If lifespan is more desirable than message throughput, the Sleep period between election cycles could be extended to add lifespan to the network, which would make HCCP have a longer lifespan for the same message throughput as LEACH.

A simulation was created with the HCCP and LEACH phases set to the same length. The results can be seen in Figures 4.13 and 4.14. LEACH has a longer lifespan, but dismal message throughput, while HCCP has excellent message throughput. The lifespan of the router mote is noticeably lower in HCCP due to HCCP leveraging the router motes to be clusterheads more frequently, which in turn is increasing the message throughput. Consequently, the router motes have shorter lifespans as the network relies on the routers to become clusterheads more frequently.

Eliminating the Clusterhead Candidacy Phase

The Clusterhead Candidacy phase is a phase in HCCP that can be considered overhead, and is an obvious target to eliminate when looking to achieve longer lifespans. The Clusterhead Candidacy phase offers very little in the way of overhead, however, since all motes that are not considering being a clusterhead are sleeping during this phase. The clusterhead motes will then have time to share information about which motes should be clusterheads for that round. If the Clusterhead Candidacy phase is dropped, the functionality must be moved to the Clusterhead Election phase.

The new Clusterhead Election phase would then work as follows:

1. Choose to be a clusterhead or not
2. Listen for clusterhead elections. If clusterhead, use Goodness Delay to discover how long to wait until it is time to Announce itself as clusterhead.
 - If a mote that has elected to be a clusterhead hears a clusterhead announcement before it sends one, choose to not be a clusterhead, and follow the mote that sent the clusterhead election.

The rest of the protocol runs as previously prescribed. The problem with overloading the Clusterhead Announcement is that when a collision occurs, it will cause much more damage to the entire cycle of the protocol. If a collision occurs during the Clusterhead Candidacy phase, one or both of the motes that are announcing their candidacy will opt to not be clusterheads, meaning that there will not be many collisions during the Clusterhead Announcement phase. When a collision happens in

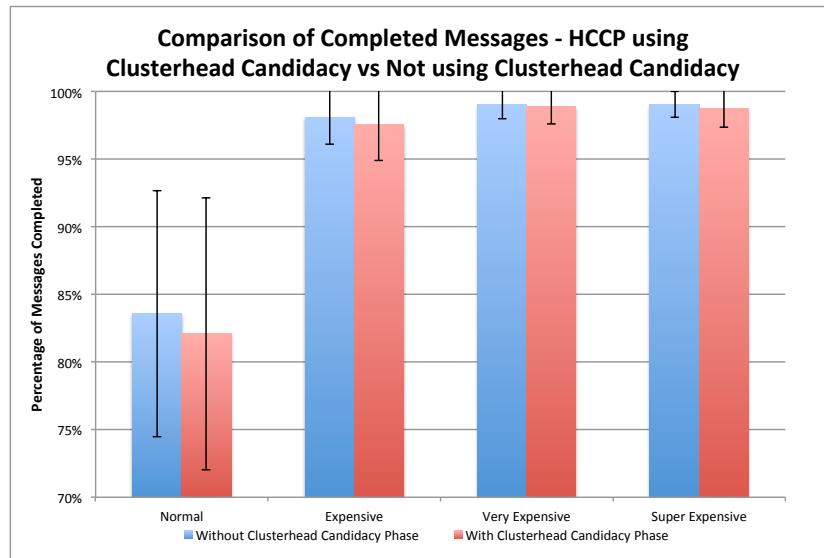


Figure 4.15: HCCP using Clusterhead Candidacy vs Not using Clusterhead Candidacy. The message throughput is comparable between the two networks.

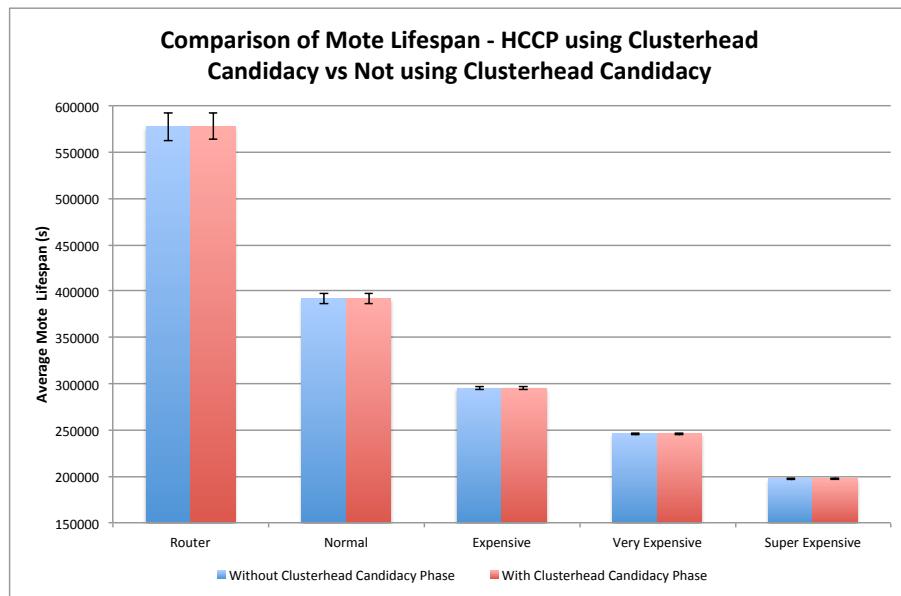


Figure 4.16: HCCP using Clusterhead Candidacy vs Not using Clusterhead Candidacy. The results are quite even in terms of mote lifespan.

the overloaded Clusterhead Announcement none of the motes in the neighbourhood will be able to follow either of the motes that made the announcements, since the messages will be garbled. This means that the motes will continue listening for clusterhead announcement messages, which means that these motes will at best be using the third best clusterhead mote in the neighbourhood (the next best mote after the two motes that had a collision).

Two simulations were created with all parameters the same, with the exception that one of the simulations used a modified version of HCCP that dropped the Clusterhead Candidacy phase. The results, as seen in Figures 4.16 and 4.15, show that whether the Clusterhead Candidacy phase is independent, or merged into the Clusterhead Election phase the results are approximately the same. This means that though merging the Clusterhead Candidacy and Clusterhead Announcement phases saves energy since motes are not on as long each cycle, clustermotes are using poorer clusterheads each round. Based on the simulation results, there is very little difference between running HCCP in either of these configurations, so either could be used in a physical deployment.

Sink Sleep vs Suboptimal Clusterheads

The HCCP Clusterhead Candidacy stage causes an issue called HCCP Blocking (previously discussed in Section 3.2.2) where messages will not flow into motes that are in range of the sink. There are two methods of allowing messages to get into the motes that are in radio range of the sink:

- allow the sink to sleep some percentage of the time, or

- have suboptimal clusterheads which will elect to be clusterhead despite the fact there is a better clusterhead in the neighbourhood.

There are drawbacks to both of these options. If the sink is asleep, no messages are being received at the sink, which could effect how many messages reach the sink. If suboptimal clusterheads are used, more motes are electing to be clusterheads, more motes are clusterheads, which puts a larger drain on more motes.

Simulations were run to see which method is better for solving the HCCP Blocking problem. Simulations were run with 20 repetitions each, using the same starting seeds. Figure 4.17 shows that there were no real differences with the percentage of messages received at the sink. This means that even if the sink sleeps every few rounds, there is no effect to the throughput of the network. This makes sense, as messages must flow into the motes adjacent to the sink before they can be hopped into the sink. Whether the sink sleeps, or motes elect to be suboptimal clusterheads, the path taken will be approximately the same. Note that router motes in the simulation did not make messages, and are therefore left off the chart.

There is also a concern that if the suboptimal clusterhead option is chosen, then the lifespan of the network might be lessened, since more motes will be clusterheads at the same time. Figure 4.18 shows us that the there is no significant hit to the network life. While difficult to see in the chart, there is no statistical difference between the two options for the ‘Super Expensive’ motes (196734 ± 611 for Sink Sleep Super Expensive motes, 196782 ± 624 for Suboptimal Super Expensive motes).

There were also no significant differences between the two methods in terms of number of collisions in the network. This is also interesting since suboptimal cluster-

heads means that there will be more traffic in more areas, clusterheads will be closer to one another.

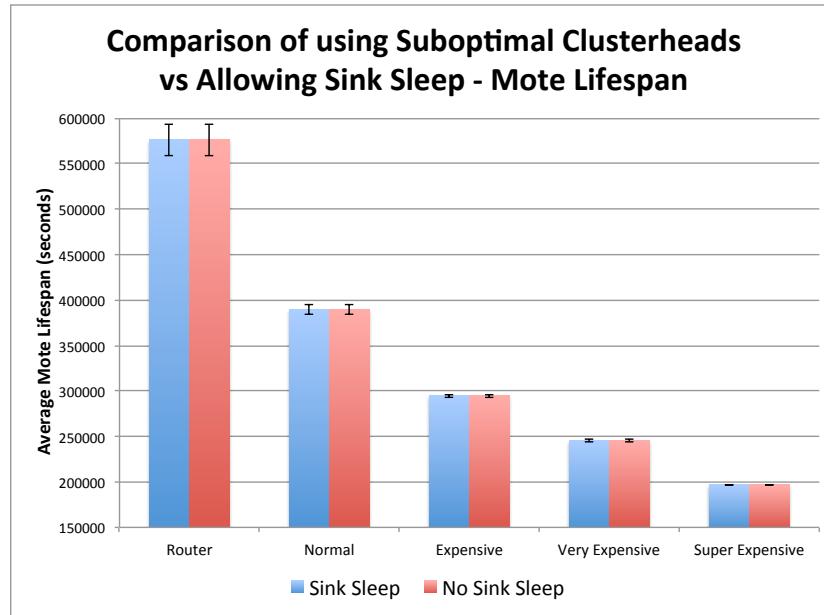


Figure 4.17: The percentage of messages created that were received at the sink is approximately the same whether the sink periodically sleeps or not.

Clusterhead Opt-out

As discussed in Section 3.1, if a network is set up to do multiple iterations of the TDMA schedule, clusterheads should be able to opt out during the Roundtable Discussion. After running some test simulations to show that this could be a successful idea, contradictions in the concept were found which showed how the idea is flawed. To successfully have a long-lived network, the Roundtable Discussion should be as short as possible, or left out altogether as the whole network must be on for the stage to be fully utilized. Since Clusterhead Opt-out requires Roundtable Discussion, its not a good fit for long-lived networks. In general, it would be better to have one longer

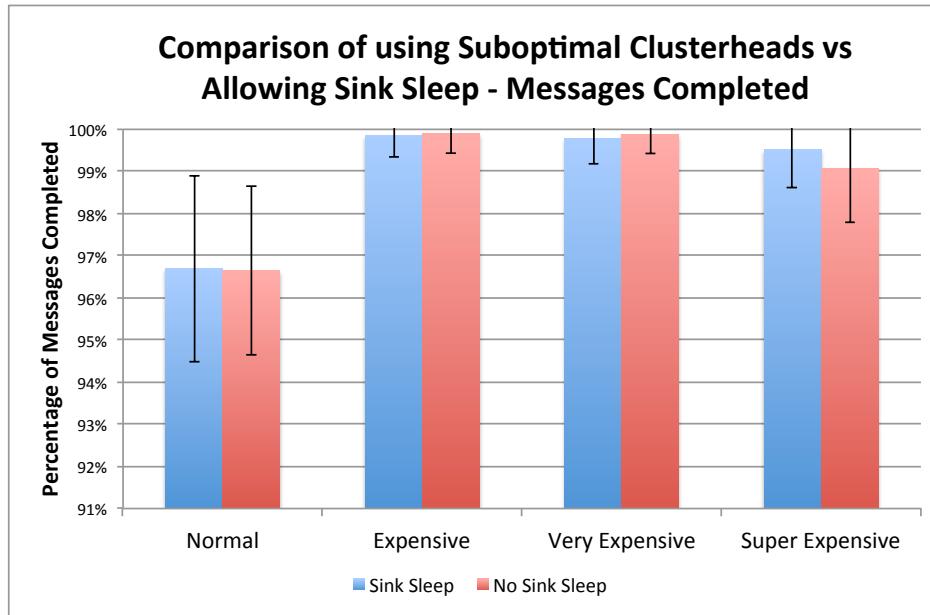


Figure 4.18: Average mote lifespan is unaffected by whether the sink periodically sleeps or not.

TDMA run schedule than a repeated TDMA schedule. A longer TDMA schedule run allows the network to send as many messages as possible, and even allows motes to turn off when they are out of messages, and turn back on and send more (provided it is still the mote's TDMA timeslice and the mote has more messages ready to send). Conversely, repeated TDMA schedules keep the same clusterhead, putting a large load onto one clusterhead and losing a large amount of network adaptivity. Another drawback is throughput. Since one mote remains a clusterhead for an extended period of time, it can't send the messages it's collecting to motes closer to the sink. This causes the messages are hopped much slower, causing message queues to fill.

There are instances where Clusterhead Opt-out could still be useful, such as networks that require quite long lifespans, where motes may go temporarily offline or die at a frequent rate. The one election with many TDMA schedules would allow the

motes to send in their readings if and when they can. If one large TMDA schedule is run and missed, motes will not be able to send in their readings since they may be dead by the time the next large TDMA schedule is published. This is, unfortunately, quite difficult to show in simulation, but presented as a possible niche case where Clusterhead Opt-out could still be useful.

Automated Ad Hoc Backbone with HCCP

There has been research into two-tier networks and networks with a backbone set of motes that are set in the network for routing messages. LEACH does not take advantage of the benefits offered by these motes, as it does not consider the heterogeneity in the network. Heinzelman et al. [21] mention that a two-tier network could be created where the routing motes collect from the neighbourhood motes, then run a LEACH schedule across the routing motes where *only* the routing motes handle routing the messages to the sink. The problem that exists with this model is that if the backbone gets broken, there is no way for the network to recover, as shown in Figure 4.19 (note that the sink is not visible in the figure). No motes past the break in the backbone will be able to send messages to the sink.

HCCP uses LEACH as a starting point, but eliminates this problem by distributing the clusterhead task to all motes in the network. HCCP will ensure that motes which are better at the clusterhead task will choose to be a clusterhead more often than motes which are not as good at the clusterhead task. This is due to the Goodness Delay that is part of the clusterhead election. Motes with larger message queues, or more battery power, or that have been given a low Sensor Mission value

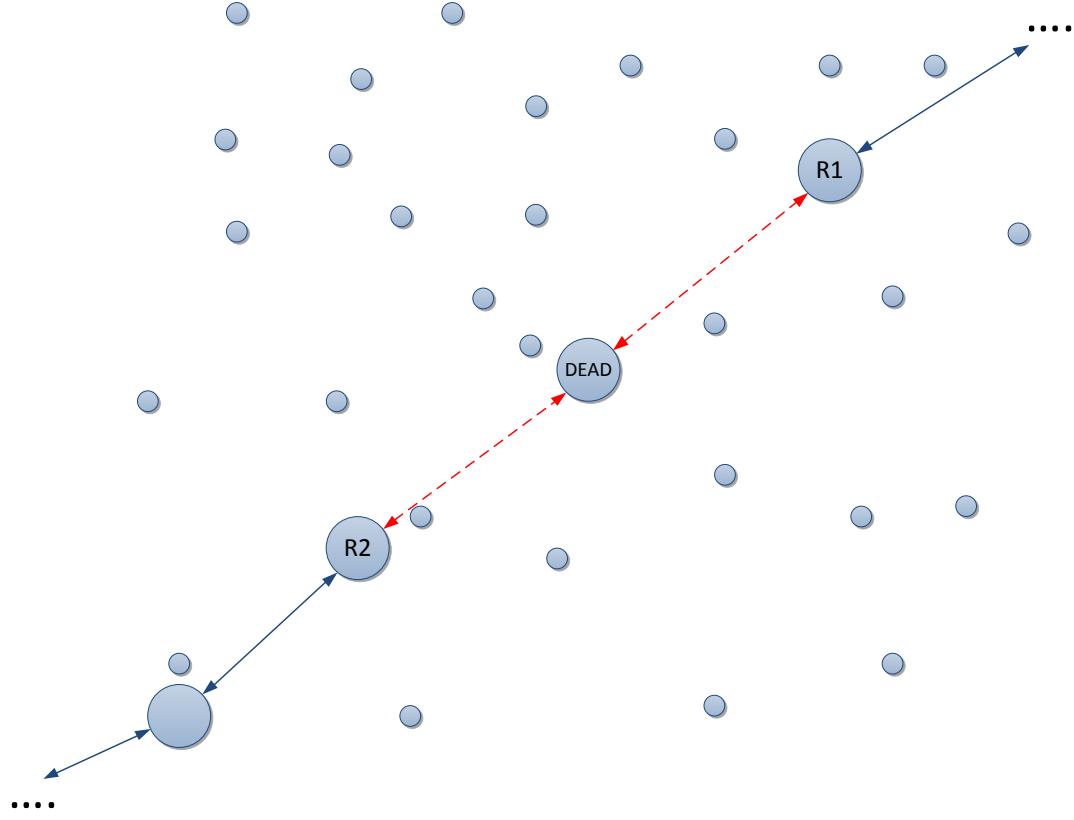


Figure 4.19: Snapshot of a part of a WSN: routing mote R1 can not communicate with routing mote R2 because the routing backbone is broken.

will choose to be clusterhead motes more often. Due to this, motes that are not good clusterheads will not choose to be clusterheads, and motes that are good clusterheads will choose to be clusterheads. This will create an ad hoc backbone across the motes which are inclined to be clusterheads. But, if the backbone is broken by a mote dying or going offline, the messages from beyond the break can still move past the break. This is because all motes can assume the clusterhead role. If there is no router mote in an area (due to the router mote going offline, or poor network distribution), other motes in the network will take on the clusterhead role. This will cause the motes that

are non-router motes to die sooner, but keep the messages from the network flowing through areas with no router motes.

The HCCP Goodness Delay creates more robust networks that can handle change within the network with minimal negative effects to the network. Messages will still be routed well due to the Roundtable Discussion sharing routing information, and messages can move over breaks in backbones since all motes can be assigned the clusterhead task.

4.2 Testbed Setup, Deployment and Results

To show the gains of leveraging heterogeneity in WSNs, a testbed deployment was created and run. Test runs using both HCCP and LEACH as the network clustering protocol were created, run and compared.

4.2.1 Building HCCP with 802.15.4

The goal of HCCP is to intelligently cluster motes in Wireless Sensor Networks, not to create a new wireless standard. HCCP works as a middleware layer between the mote's application and the wireless radio that is using some wireless protocol, such as Bluetooth, ANT radios, or even WIFI (802.11). The network protocol used in the testbed deployment was 802.15.4.

Building on top of 802.15.4 requires that the protocol's requirements are met. A basic packet in 802.15.4 is constructed using several segments that have information in well-defined places, as seen in Table 4.1. The physical layer header preamble is handled by the hardware and is outside the scope of this work. It is assumed that the preamble is functioning properly, and no changes need to be made to it for the support of HCCP.

The bytes following the physical layer preamble are the Media Access Control (MAC) layer header bytes. The bytes in the header are well-defined by the 802.15.4 protocol, with each byte given a special meaning in the header. The first two bytes are the Frame Control Field (FCF), which can be seen in Table 4.1, that tell the packet parser what to expect in the rest of the packet. The FCF details values which affect how long the packet header will be such as whether the security is enabled, and

which type of MAC addresses the packet header contains.

The next byte is the packet sequence number, which is a value that gets incremented for every packet the sending mote sends. It is used to recreate data that is split across multiple packets that are being sent to the receiving mote. In the testbed deployment, data is assumed to be contained within one packet. The sequence number of the packets is therefore irrelevant, but can not be left out without violating the guidelines described by the 802.15.4 standard.

Octets: variable	2	1	4-20	0-14	n	2
Preamble Sequence	Frame Control	Sequence Number	Addressing Fields	Security Header	Data Payload	Checksum

Table 4.1: Breakdown of data packets in 802.15.4. Adapted from [3].

Frame Control	Sequence Number	Destination Address	Source Address	Auxiliary Security Header	HCCP Packet Type	HCCP Packet Payload	Checksum
------------------	--------------------	------------------------	-------------------	---------------------------------	------------------------	------------------------	----------

Table 4.2: Customized 802.15.4 packets for use in HCCP with more detail.

Following the sequence number is the addressing field, which can be 4-20 bytes. The range is due to the way that 802.15.4 addresses can be specified. 802.15.4 has 20 byte addresses that are designed to be globally unique identifiers, but can also use 4 byte addresses to create lightweight packets. When using the 4 byte addresses there is a good chance that two motes in the network could have the same address. This causes the tradeoff of reliability and understandably of data versus the per packet cost to send a message in the network. There is both a source and destination address specified in the addressing field.

An optional security field follows the sequence number. This field could be used to encrypt the data in the packet, but it is not used in the testbed deployment. HCCP

can be used with encryption using this field with little or no modifications to the HCCP protocol.

Finally, the data payload for an 804.15.4 packet is where all the HCCP logic is added. HCCP divides up the data section of the packet into more well-known fields, much like the creators of 802.15.4 have done. HCCP divides the data section in to 8 bits for packet type, and the remainder is the packet data. The custom HCCP packet format is shown in Table 4.2. The The HCCP values for the HCCP Packet Type field are as follows:

1. **Clusterhead Candidacy Announcement** - 0x0 - The sender of this message is going to be a clusterhead candidate.
2. **Clusterhead Announcements** - 0x1 - The sender of this message is a clusterhead. Following the HCCP header is the mote's ID.
3. **Join Cluster** - 0x2 - The sender of this message wants to join the cluster of the recipient. The mote that has the address in the addressing field is the mote it will follow. The clusterhead that is the recipient will use the 'from' address as the ID of its new clustermote.
4. **TDMA Schedule** - 0x3 - A broadcast message to all surrounding motes with the TDMA schedule the clustermotes will be following. A sample TDMA schedule can be seen in Table 4.3.
5. **Data Packet** - 0x4 - Messages sent from clustermotes to their clusterhead. n data bytes will follow the type field.

Packet Type	Runtime for each Mote (s or ms)	First ID & Delay	Second ID & Delay	...	nth ID & Delay
0x3	20	12:20	18:40	...	88:70

Table 4.3: Sample TDMA schedule as sent by the clusterhead to its clustermotes.

6. **Roundtable Discussion Packet** - There are a few well-defined packets that can be sent during a Roundtable Discussion. To save in the packets, they can share the type field with the other well-defined packets, or can be assigned type values of their own.
- (a) **Clusterhead Opt-out** - 0x5 - Clusterhead can quit if there are multiple iterations of the HCCP schedule.
 - (b) **Roundtable Clusterhead Announcement** - 0x6 - a replacement clusterhead if an existing clusterhead opts out.
7. **Other** - 0x8 - If any packets are defined for a customization of HCCP, this value can be used. An instance where other packets could be used is during the Roundtable Discussion, if the motes need to share some other information. In the testbed deployment, routing information is shared using the other packet type during the Roundtable Discussion time. The contents following the HCCP header are at the discretion of the network administrator.

After the HCCP packet type field is the HCCP packet data field which contains the packet payload. After that the 802.15.4 checksum follows to make the packet 802.15.4 compliant.

4.2.2 Description of Testbed Motes

To create a heterogeneous WSN, motes from multiple vendors were selected for the experimental WSNs. The motes selected have different radios, run different operating systems and have differing capabilities. Below is a description of the motes used for the experiments, the capabilities of the motes, the modifications that could be possible to increase heterogeneity and the modifications made to them.

Sun Microsystems SunSPOT

Sun Microsystems Labs [35] saw the need for easily programmable motes with solid form-factor and simple tools to use. SunSPOT stands for ‘Sun Microsystems Small Programmable Object Technology’, which is designed to be simple to both program and deploy. Sun had already created a version of Java to run on mobile phones called Java 2 Micro Edition (J2ME) [36]. Sun used this code base for a starting point for the SunSPOT programming interface.

The SunSPOTs are sold in boxes with 2 motes called “free-range” motes seen in Figure 4.20, and one base station mote, as seen in Figure 4.21. The free range motes are complete with a rechargeable battery that charges from USB, a CC2420 radio, a number of sensors and LEDs for the programmer to use. The free range SunSPOTs are divided into 2 boards, as seen in Figure 4.22. The top board is the board visible on the top of the SunSPOT, which is the sensor board that holds all the sensors, buttons and LEDs. The second board is the main board that hosts the microcontroller and the radio. The base station SunSPOTs share the same main board as the free range SunSPOTs, but do not have a rechargeable battery or a sensor board. The common

main board is convenient, as the radio code can be shared between the both the free-range motes, and the base station.



Figure 4.20: The Sun Microsystems SunSPOT.



Figure 4.21: A SunSPOT free-range mote and SunSPOT base station.



Figure 4.22: A SunSPOT taken apart to view the components in the casing.

Possible Modifications to the SunSPOTS. The rechargeable battery the SunSPOTS ship with is small and does not hold enough energy to power the SunSPOT for a long period of time. A simple modification is to remove the back of the SunSPOT, replacing the battery with any other power source providing up to 4.9 volts as stated in the SunSPOT developer's manual [37]. This could be done with 3 AA batteries, which would provide much more available energy.

Uses for the SunSPOT in a Heterogeneous WSN. The base station the SunSPOTS ship with makes a very convenient base station for a WSN. The radio code developed for the free range SunSPOTS can be used on the base station while it is connected to a computer. The base station is designed to run with Java code developed using the SunSPOT suite developed for the SunSPOT, and the base station can be easily accessed via a Java program. The program running the base station can

then log the data to a database, or post it to a website to make the data collected by the WSN available world-wide immediately.

The base station could also be used to make an entire computer a mote. That is to say, the base station could be connected to a computer but not act as a sink. This would allow the mote to have nearly unlimited power and memory resources, making it an ideal routing mote.

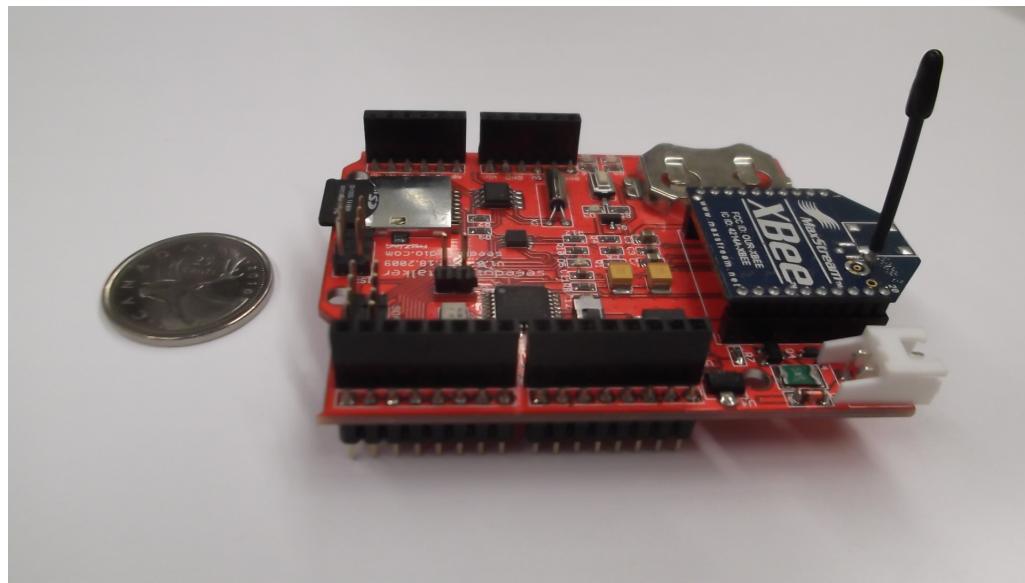


Figure 4.23: Seeeduino Stalker with XBee port exposed.

Seeeduino Stalker and Arduino Technology

Seeed Studios [38] is a developer and retailer of microcontrollers online, specializing in Arduino-compatible microcontrollers. One of the products they provide is a mote called the Seeeduino Stalker. The Stalker is a microcontroller based on the Arduino design [39], using a compatible processor and standard Arduino pin layouts. As seen in Figure 4.23 the Stalker has an onboard battery, and logging capabilities,

using MicroSD cards to log to. The Stalker does not ship with a built-in radio, but makes an XBee compatible slot available. To allow the Stalker to communicate with the other motes, we chose the XBee Series 1 modules - which communicate using the 802.15.4 standard [3].

The Stalker does not have any native sensing abilities, but can be configured to use either Arduino-compatible shields (which attach to the female pin headers on the board to provide some service). This makes the Stalker ideal for creating highly customized motes. The Stalker could be configured to have no sensors to make a low-power routing mote with the ability to log large amounts of data on its MicroSD card. Or, can be configured to use a variety of sensors and be a specialized sensor mote.

A problem does arise with the Stalker's battery power. The Stalker is sold with a port to use a 'button battery', as seen in Figure 4.23. More power is very desirable, so the Seeeduino Stalker motes will need to be fitted with a larger battery supply before it could be used as a useful mote. The Stalker is sold with a port to plug in an external power supply, so this is not a large problem.

A major problem with the Stalker is the CPU it uses. The CPU has very limited program space. The code developed frequently overran the size of the program space, causing the mote to crash. Due to this, a minimal installation of Contiki OS has been created to allow the Stalker to work in the network with limited capabilities. The Stalker can only participate in the network as a clustermote, not as a Clusterhead.

Another Arduino mote was made for the network. The Arduino Duemilanove is a more powerful Arduino with more memory and fewer built-on features. The

Duemilanove required an extension called a shield to allow the board to use a radio, as seen in Figures 4.24 and 4.25. By adding a radio extension, the Arduino gains the ability to send data over an 802.15.4 network, but does not replicate the logging ability the Stalker has.

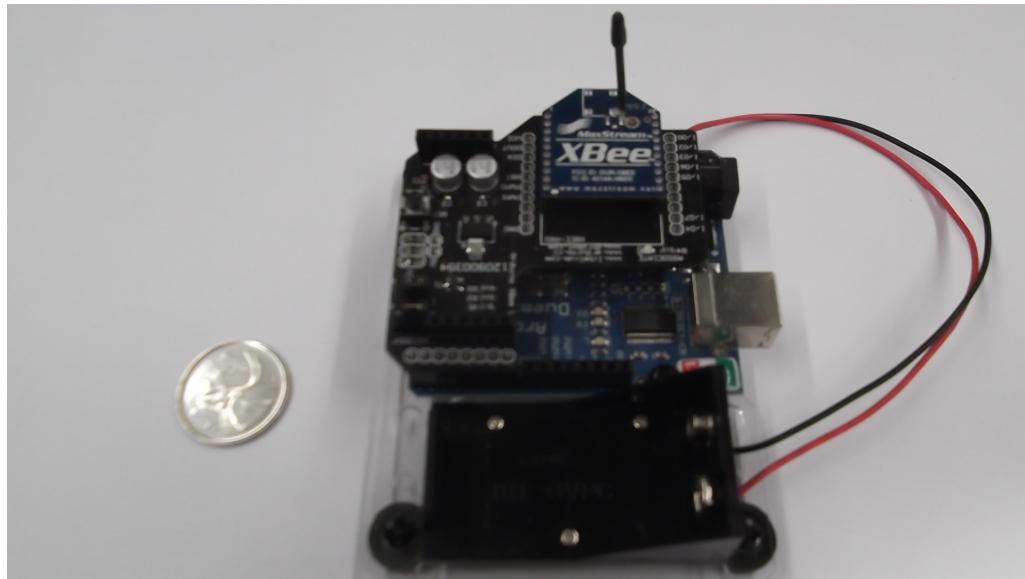


Figure 4.24: An Arduino Duemilanove turned into a mote using the XBee Shield and an XBee radio.

Possible modifications to the Stalker. There is no battery packaged with the Arduino Duemilanove, so adding a battery is necessary for the mote to work in the network. The Duemilanove is a custom setup, and could be considered to be a heterogeneous mote by design.

The Stalker ships with a button-cell battery that would not provide the stalker with a very long lifespan. A larger battery would be necessary to make the Stalker a useful mote.

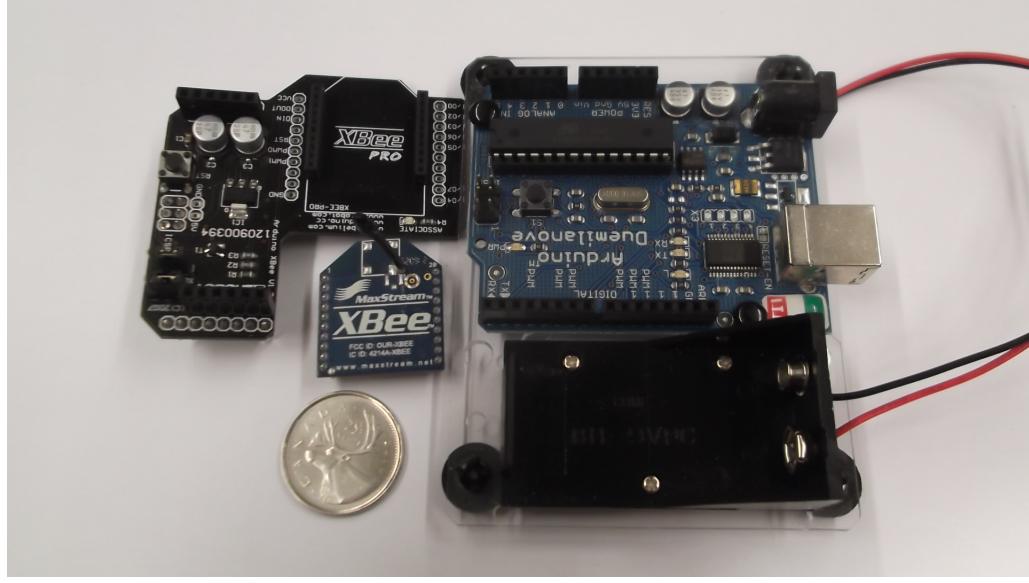


Figure 4.25: The parts used to make the Arduino mote.

Uses for the Stalker and Arduino in a Heterogeneous WSN. The Stalker could be used for a routing mote, utilizing it's large storage capabilities provided by the MicroSD card. The Stalker could also be used as a customized sensing node, using an Arduino-compatible shield, or by using the Stalker's general purpose pins to connect to a sensor.

Software used. Since the Stalker does not have a radio built onto the board, a third-party radio must be used. The Stalker provides an XBee-compatible port to host a radio. This XBee-compatible port uses the serial port on the Stalker board to contact the radio. This works well with XBee radios, as they natively work in ‘transparent mode’, which means that the serial transmissions are sent over-the-air to other XBee radios in a way that the software using the radio need not know that a radio is being used.

Since this project requires developing custom packets, a different mode must be used when using the XBee radios to send and receive custom packets. The XBee Series 2, which is used in the experiment, has a mode called ‘API mode’, that allows a developer to create and read packets manually. A library named Arduino-Xbee [40] was used to create and parse 802.15.4 packets properly. The library was modified to be more flexible and configurable to be more useful in the experiments.

tMote Sky

The origins of the tMote Sky are from University of California, Berkeley. The basic design is named Telos [41], and has been marketed by several companies. The Telos mote was marketed by MoteIV, renaming it the tMote Sky. MoteIV has been rebranded and renamed Sentilla. Sentilla is still marketing the motes, as the Sentilla Mica. The motes have also been marketed as Crossbow, renamed to JCreate, all of which have a similar design as seen in Figure 4.26. Each of these rebranded Telos motes have slightly different processors with slightly differing memory sizes and sensors. The Telos variants are therefore heterogeneous due to all of the different variants that are available.

Possible modifications to the tMote Sky. The Telos mote has a large built-on power supply, and should not need a larger power supply. The Telos does not have a large onboard memory, adding an external memory would make it a good router/clusterhead mote.



Figure 4.26: The tmote Sky, which is a Telos variant.

Uses for the tMote Sky in a Heterogeneous WSN. There are a number of onboard sensors making this a good mote for sensing tasks. Though, the large battery also makes the tMote a good candidate for clusterhead, making the tMote a very versatile mote.

4.2.3 Heterogeneous Setup of the Motes

Table 4.4 summarizes the motes, comparing the processors, radios, memory and battery power. From the table it is easy to see the stark differences between the motes. The SunSPOT has a huge amount of memory and much more powerful processor than the other motes, but its battery is tiny. The tMote has a full operating system that is very easy to work with, and a large battery, but not a lot of flash memory. Though, the amount of flash memory is large in comparison with the Stalker and the Arduino which have almost no memory and very small extended memory.

Clearly, the SunSPOT is the most powerful mote in every respect except battery power. The tMote will take the clusterhead head task the most frequently due to it's good extended memory size and battery power. The Stalker and Arduino will not likely be clusterheads due to the limited extended memory which equates to limited queue sizes available on the devices.

While creating the testbed deployment, developing the software for LEACH and HCCP on the Stalker was an issue, due to the very limited program space. Due to this, a minimalistic, chopped down version of Contiki OS had to be made. The software uses the message queue code and interrupt callback code from Contiki, while removing all other functions of the OS. Contiki functions that were removed were any threads and context switching, drivers for sensors and portability code so this version of Contiki can only run on Arduino microcontrollers. This code was ported to the Arduino with the ATmega 328 chip first, then an even more stripped-down version had to be made for the Stalker with the less powerful ATmega 168p chip. The Stalker version of the software turns the message queue into a ring buffer, and has limited roundtable abilities which only allow the Stalker to accept routing information, and do nothing else during the Roundtable time.

Clearly, any of the modifications previously discussed could add more heterogeneity to this already highly heterogeneous network if so desired.

4.2.4 Testbed deployment

To test the real-world feasibility of HCCP, a controlled physical deployment was run. SunSPOTs were used exclusively in the controlled deployment, as they were

	Processor	Radio	Memory	Extended Memory	OS / Software
SunSPOT	400 MHz 32-bit ARM920T	CC2420	1 MB	8 MB	SunSPOT Framework
tMote	8 MHz 16-bit TI MSP430	CC2420	10 kB	48 kB	Contiki OS
Stalker	20 MHz 8-bit Atmel ATmega 168P	XBee Socket	1 kB	16 kB	Minimal Contiki OS
Arduino	16 MHz 8-bit Atmel ATmega 328	XBee Socket	2 kB	32 kB	Minimal Contiki OS

Table 4.4: Comparison of Motes in Testbed Deployment.

readily available in large quantities. A problem that plagued both testbed deployments was that the SunSPOTs used were about five years old, and the rechargeable batteries that are built-in were starting to show signs of age. In the deployment, some of the motes had the low battery warning light illuminated immediately despite being fully charged. Due to the poor batteries, some of the motes started dying after only 4 hours of running the experiment. Due to the wear on the rechargeable batteries, all the motes had a different amount of charge to use, which made the network heterogeneous in terms of battery power. Adding different types of motes would make the network more heterogenous, but the results would be very difficult to interpret due to the amount of heterogeneity.

LEACH was set up as prescribed by Heinzelman et al. [21] adding beacon routing information to the Clusterhead Election announcement.

House Monitoring Deployment

To show HCCP working in a real-world scenario, motes were distributed around a house. Setting out motes around a house can simulate house monitoring for temperature, humidity or even security. Motes could also be used in a house for home automation such as turning on lights in the room as a person enters, or having moving what's on television from one room to the next as a person moves through a house.



Figure 4.27: Motes distributed around a house.

The motes were deployed at various heights across a house as seen in Figure 4.27, with some motes placed outside in waterproof containers. The networks had 44 motes and 1 sink distributed over (110) square meters (about 1200 square feet) on various different levels and were made to communicate through various different materials. The motes relayed their messages back to a sink that was placed centrally in the house, this provided many routes back to the sink from any given mote. The motes were setup to create five new messages per minute, with all messages to be routed towards the sink. Radio range was reduced to 4 meters line-of-sight to keep the

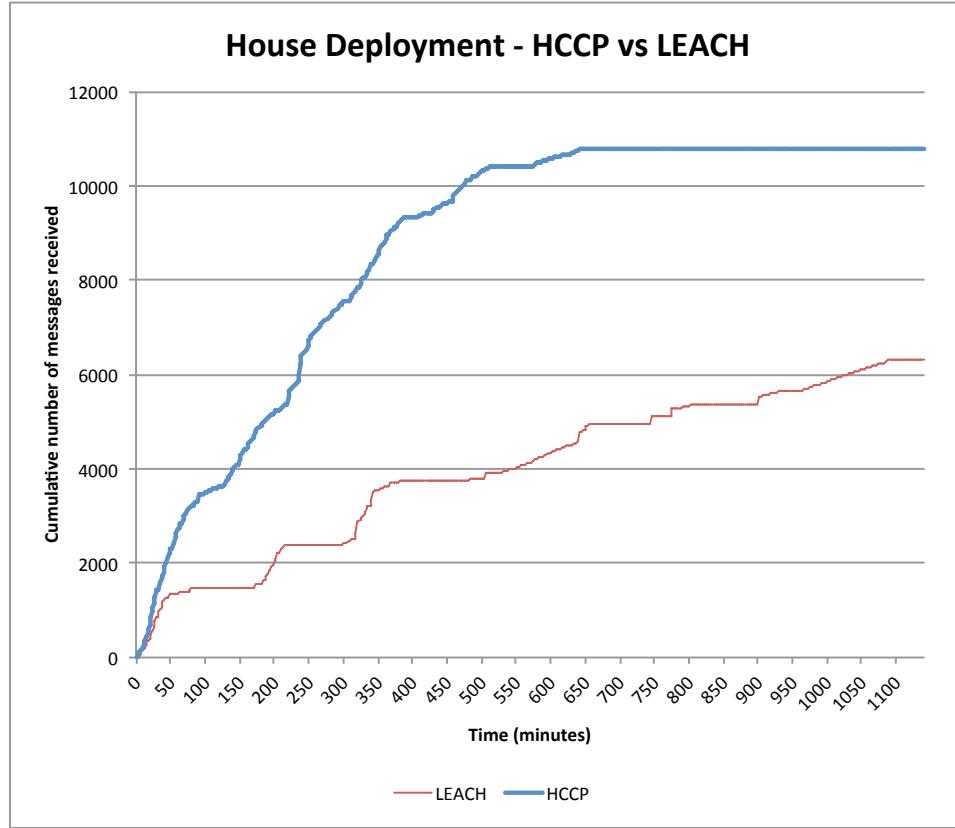


Figure 4.28: Cumulative number of messages received from the deployment done in a house.

network maintainable and to ensure the network was a multi-hop network. Motes on the edges of the deployment were setup to be at least three hops away. The same network layout was used with both HCCP and LEACH deployments.

HCCP proved to be very effective at moving messages to the sink faster, as seen in Figure 4.28. At 600 minutes (10 hours) the motes closest to the sink became inoperable, stopping the progress of the network. Since the network was setup as quite sparsely, a few motes dying early near the sink would have a large effect on the network. The performance of LEACH also degrades after the 600 minute mark, having received 62% of its messages before the 10 hour mark.

Before the 10 hour mark, HCCP achieved approximately double the message throughput that LEACH did. Also, HCCP had a much more consistent curve, creating a more linear chart. LEACH has a step-function, probably due to the 5% clusterhead rate as prescribed by Heinzelman et al. in their description of LEACH. Since the network cycle was setup to take one minute, messages can only flow into the motes next to the sink every 20 minutes, when the motes next to the sink choose to become clusterheads.

HCCP, on the other hand, was focused on Message Queue (95%) and Battery Power (5%). As the queues of the motes filled, they would be less likely to be clusterheads, allowing messages to flow out of them. Then, as motes further away from the sink get full of messages, they are less likely to become clusterheads. As the motes further away choose to be clusterheads less frequently, the motes close to the sink will more likely be clusterheads, which allows the messages flow towards the sink. As the motes close to the sink fill with messages from the motes further from the sink fill with messages, they will not be clusterheads as frequently, which allows the messages to be sent to the sink.

Grid Deployment

A second deployment was done with 30 motes and a sink. The motes were set up in a 5 by 6 grid, with 1.5 meter spacings between them. This made the grid spread out over a 7.5 meter by 9 meter area. The radio range was set to a 6 meter range, which will degrade over the lifespan of the network. The sink was set up at the side of the network, as seen in Figure 4.29.

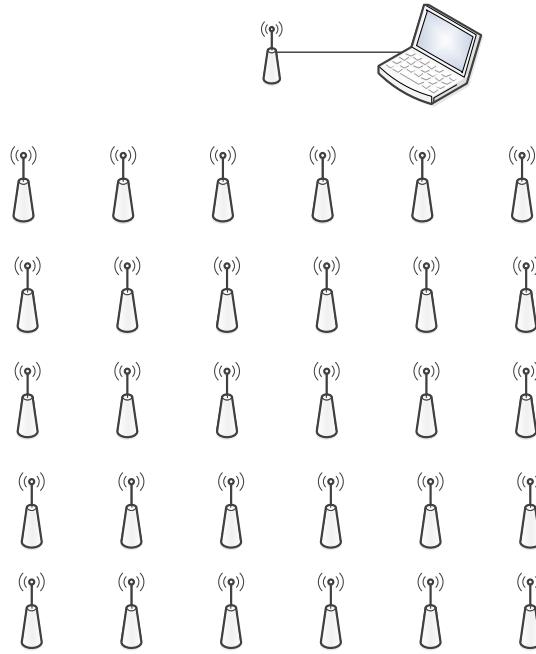


Figure 4.29: The grid deployment setup.

During both the LEACH and HCCP deployments, some of the motes immediately had the battery low warning light illuminated, despite the fact all batteries were fully charged before the deployment. Rechargeable batteries lose their ability to hold a charge as they age; since these are older motes it is only natural that some of the batteries would be quite degraded. The motes that were nearly dead at time of deployment were left in the network for the test.

One problem about battery-powered devices is that as the battery drains, the device becomes unreliable. This unreliability stems from transistors not switching properly due to not being able to draw enough current, or not having enough voltage to cause the transistor to change state. Since the SunSPOT batteries were old, the motes were only reliable for about 10 hours. After 10 hours the motes could not maintain the network schedule, as the timer would not work properly. Interestingly,

during the HCCP deployment, 3 motes lasted over 60 hours but didn't send any messages during that time, meaning that the motes were in deep sleep for that entire time. Not waking from sleep can happen when the wakeup timer never fires, leaving the mote asleep indefinitely.

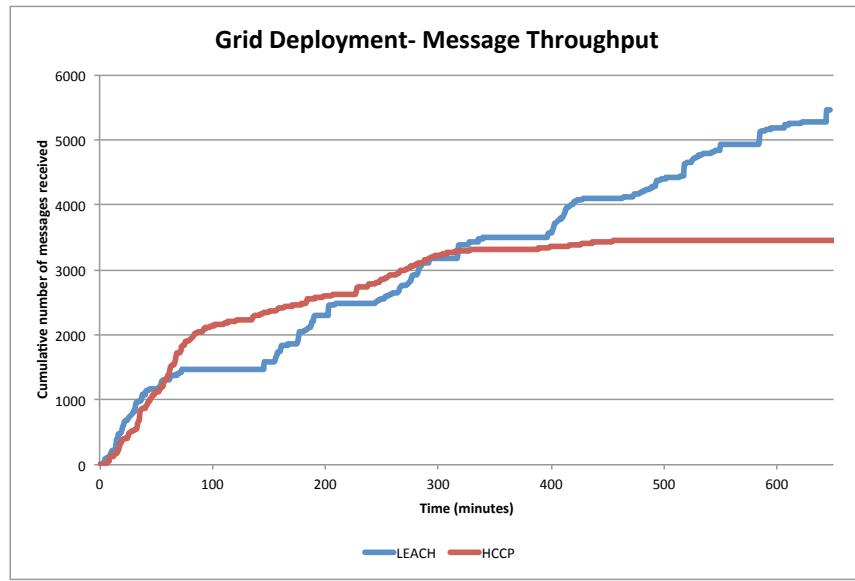


Figure 4.30: Results from LEACH and HCCP Grid Deployments

The results from the grid deployment is shown in Figure 4.30. At the start of the deployments, HCCP started with a better message throughput, as expected from the results of the simulations. The throughput of HCCP then drops to a slower rate, allowing LEACH to catch up, and overtake it. HCCP then levels off, as motes begin to fail as the 300 minute (5 hour) point approaches. Motes close to the sink were failing due to going to sleep, causing the messages to not have a route to the sink.

The LEACH curve is more consistent than the HCCP curve, having steady steps. The steps are quite interesting, as the LEACH network cycle was setup to take a minute, so the steps should be in shorter cycles than 10 minutes.

If the motes had been provisioned with larger batteries, HCCP should maintain the slope it was forming before the 100 minute mark. The HCCP deployment had the unfortunate luck of having the motes closest to the sink die first due to poor batteries, dying early of reasons unrelated to the protocol.

For network lifespan, the last message received in the LEACH network was at 646 minutes (10.75 hours). Since the motes closest to the sink in the HCCP deployment had problems waking from sleep, the last message was received at 3489 minutes (58.15 hours). This huge network lifespan is due to a mote rebooting after being in deep sleep for a long time, rejoining the network and sending its messages to the sink. LEACH likely had motes do the same thing, but the motes weren't as close to the sink and therefore did not create these odd data points.

Overall, HCCP has worked as a proof-of-concept in a real-world deployment. Since older motes were used, the results are difficult to interpret. Despite these difficulties, HCCP worked well, showing its real-world plausibility.

Chapter 5

Conclusion and Future Work

In both simulation and testbed deployments, HCCP has increased message throughput while having very little negative effects on network lifespan. HCCP provides a robust and tuneable network that provides many options to elect motes that have certain properties as clusterheads more frequently. Depending on the configuration, HCCP appears to be able to provide network lifespans as long as LEACH while having almost double the message throughput, as shown in both simulations and physical deployments.

While HCCP is more complex than LEACH, it has a more robust design, and provides infrastructure for disseminating information across the network. HCCP relies on the network to have schedules in tight synchronization, which is a requirement in all WSN protocols, the novelty that HCCP adds is to use this highly coupled time to describe how good a mote would be at being a clusterhead. The Goodness Delay utilizes time that would otherwise not be used to optimize which motes will be clusterheads.

HCCP is designed to be an easily tuneable protocol, giving the network designer tools to make the best network possible. An example of tuning the network is changing the length of the Roundtable Discussion time. This time is allotted to share any network-specific information that doesn't fit anywhere else into the network. HCCP provides some suggestions as to what to share during this time, but is intended to be customized for every individual network, or even left out if desired.

Heterogeneity in a WSN can be a powerful thing, and HCCP uses this heterogeneity to elect better clusterheads. The definition of 'better' is up to the network designer to choose. HCCP allows the election to focus on motes that have qualities that a network designer feels are important to the network. This thesis provides insight into what heterogeneous factors have to offer to the network, how they affect the election, message throughput and network/mote lifespans.

Overall, heterogeneity, whether the differences are small or large, can be used to generate better functioning WSNs. Choosing motes that are better for the task of clusterhead will prevent lost messages with no cost to the lifespan of the network. The gains of using heterogeneity that is already inherent in every WSN is free, having few negative effects. Having the heterogeneity of a network self-assessed will ensure that as a network degrades over time that the best possible motes at that time will be doing tasks that they are well-suited for, providing gains to the network from the time it is deployed to the time the last mote ceases to function.

Looking forward, HCCP could be simplified for further energy savings. For instance, the Clusterhead Candidacy phase could likely be merged with the Clusterhead Announcement phase, as there is little difference to the network lifespan or message

throughput when the phases are merged. Since removing the Clusterhead Candidacy would make HCCP simpler and would not have any negative affects on the network, it would probably be a good idea. Further research could be done on how HCCP could surpass LEACH in terms of network lifespan.

A better routing protocol should be incorporated into the design of HCCP. More research could be done to see which routing algorithms would be the best with HCCP. The routing algorithm should also take advantage of the heterogeneity of the network, routing messages through stronger motes. HCCP was designed with this in mind, and is easily changed to use any routing algorithm. Further, if a mobile sink is a consideration, then using a routing algorithm that works efficiently with networks that have a mobile sink, such as MobiRoute [42], is also possible.

Querying motes in the network from the sink was not a focus of HCCP. If a network administrator wants to request data from a mote, there is no set method of sending the messages to any mote in the network but the sink. To add querying to HCCP, a different routing method would need to be used, as only the sink's position is advertised while using beacon routing. A routing protocol that has a table of all the mote's last known positions would need to be used. Once routing is set up, the querying framework would need to be added. This feature could be added into the Roundtable Discussion time, or into a new phase that is dedicated to mote querying.

HCCP was not designed with any inherent security measures. If messages are solely to be routed to the sink, simple RSA encryption [43] could be used to encrypt the messages that are being sent to the sink. Other security issues such as an errant mote that is dying continually becoming a clusterhead could be solved by blacklisting

motes that are continuously clusterheads.

To create a WSN with fewer collisions in the network, it is possible to use Frequency Division Multiple access (FDMA) [44] to use different radio channels for different clusters. Clusterheads could announce the channel along with the TDMA schedule. Using multiple radio channels during the TDMA clustermote reporting run time would remove many of the collisions that happen during the TDMA run time.

Radio power could also be adjusted to suit the size of the network. Adjusting the radio power to the lowest transmission power saves mote energy, and prevents motes from flooding the network with unnecessary messages from motes. Also, smaller transmission ranges make clusters smaller. Changing the radio power effectively changes the cluster size, since few motes will be in range. Lin et al. [45], Zheng et al. [46], Xiao and Yu [47] and others have developed efficient methods of adjusting the radio power to enable motes to talk to a limited number of neighbouring motes. Messages to communicate radio power could be added to the HCCP's Roundtable Discussion, or piggybacked on announcement messages, such as clusterhead announcements.

5.1 Future Goals of WSNs

Smart Dust [48, 49], is one of the potential futures for WSNs. As circuits, processors and motes get smaller and smaller there may come a point where motes are no larger than the size of dust. These motes would have near negligible unit cost and be deployed in the thousands. Smart Dust would communicate wirelessly, monitoring some phenomena and glean energy from the environment to stay powered. Smart Dust would need to be deployed with a dense network that would need to be

a multi-hop network, as each device would have minimal battery power and minimal transmission power. Any protocols designed would have to be scalable to accommodate large networks of Smart Dust. HCCP has been designed to scale, but smart dust might be a larger scale than HCCP was designed for, though HCCP provides interesting ways to use the differences in the dust well.

As devices get smaller and more feature-rich, the eventual outcome might be that all devices are networked and can communicate with each other. When all devices can communicate with each other, they would create ‘an internet of things’, a term used frequently by Adam Dunkels [50], who is the creator of Contiki OS [51, 52] and a key member of the WSN community. Dunkels et al. [53] believe that devices will communicate with each other using a lightweight version of IPv6 called 6LoWPAN IPv6 [54]. A subset of the larger IPv6 protocol [55] is an obvious choice, as it can provide unique addresses for 3.4×10^{38} devices, which would provide 48 000 000 000 000 000 000 000 000 000 addresses for every one of the 7 billion humans on Earth. This means that every single device in the Internet of Things could have its own address, and be able to communicate with any other device in the internet of things. This would make the entire Internet a massive scale WSN of sorts, with many heterogeneous devices all connected using the common language of IPv6. HCCP uses IPv6 addressing, allowing networks running HCCP to possibly be part of the internet of things.

Bibliography

- [1] Ian F. Akyildiz, Weilian Su, Yogesh Sankarasubramaniam, and Erdal Cayirci. A survey on sensor networks. *IEEE Communications Magazine*, 40(8):102–114, Aug 2002.
- [2] Arati Manjehwar and Dharma Agrawal. TEEN: A routing protocol for enhanced efficiency in wireless sensor networks. In *Workshop on Parallel and Distributed Computing Issues in Wireless Networks and Mobile Computing*, pages 2009–2015, April 2001.
- [3] Institute of Electrical and Electronics Engineers (IEEE). The 802.15.4 protocol. <http://www.ieee802.org/15/pub/TG4e.html>, Accessed on February 1, 2011.
- [4] Deborah Estrin, Ramesh Govindan, John Heidemann, and Satish Kumar. Next century challenges: Scalable coordination in sensor networks. In *ACM/IEEE International Conference on Mobile Computing and Networking*, pages 263–270, August 1999.
- [5] Alan Mainwaring, David Culler, Joseph Polastre, Robert Szewczyk, and John Anderson. Wireless sensor networks for habitat monitoring. In *Proceedings of the*

- 1st ACM International Workshop on Wireless Sensor Networks and Applications*, pages 88–97, September 2002.
- [6] Pei Zhang, Christopher Sadler, Stephen Lyon, and Magaret Martonosi. Hardware design experiences in zebranet. In *Proceedings of the 2nd International Conference on Embedded Networked Sensor Systems*, pages 227–238, November 2004.
- [7] Ning Xu, Sumit Rangwala, Krishna Kant Chintalapudi, Deepak Ganesan, Alan Broad, Ramesh Govindan, and Deborah Estrin. A wireless sensor network for structural monitoring. In *Proceedings of the 2nd International Conference on Embedded Networked Sensor Systems*, pages 13–24, November 2004.
- [8] Geoffrey Werner-Allen, Konrad Lorincz, Matt Welsh, Omar Marcillo, Jeff Johnson, Mario Ruiz, and Jonathan Lees. Deploying a Wireless Sensor Network on an Active Volcano. *IEEE Internet Computing*, 10:18–25, March/April 2006.
- [9] Omprakash Gnawali, Rodrigo Fonseca, Kyle Jamieson, David Moss, and Philip Levis. Collection Tree Protocol. In *Proceedings of the 7th ACM Conference on Embedded Networked Sensor Systems*, pages 1–14, November 2009.
- [10] Nikola Tesla. *Experiments with Alternate Currents of High Potential and High Frequency*. Gutenberg Press, 2004.
- [11] Gregory Pottie and William Kaiser. Wireless integrated network sensors. *Communications of the ACM*, 43:51–58, May 2000.
- [12] Leonard Kleinrock and Fouad Tobagi. Packet switching in radio channels: Part

- i-carrier sense multiple-access modes and their throughput-delay characteristics. *IEEE Transactions on Communications*, 23(12):1400–1416, December 1975.
- [13] Fouad Tobagi and Leonard Kleinrock. Packet switching in radio channels: Part ii—the hidden terminal problem in carrier sense multiple-access and the busy-tone solution. *IEEE Transactions on Communications*, 23(12):1417–1433, December 1975.
- [14] Alec Woo and David Culler. A transmission control scheme for media access in sensor networks. In *Proceedings of the 7th annual international conference on Mobile computing and networking*, International Conference on Mobile Computing and Networking, pages 221–235, July 2001.
- [15] Sinem Coleri Ergen and Pravin Varaiya. TDMA scheduling algorithms for wireless sensor networks. *Wireless Networks*, 16:985–997, May 2010.
- [16] Curt Schurgers, Vlasios Tsiatsis, Saurabh Ganeriwal, and Mani Srivastava. Optimizing sensor networks in the energy-latency-density design space. *IEEE Transactions on Mobile Computing*, 1(1):70–80, January 2002.
- [17] Chalermek Intanagonwiwat, Ramesh Govindan, and Deborah Estrin. Directed diffusion: A scalable and robust communication paradigm for sensor networks. In *ACM International Conference on Mobile Computing and Networking*, pages 56–67. ACM, 2000.
- [18] Sandra M. Hedetniemi, Stephen T. Hedetniemi, and Arthur L. Liestman. A

- survey of gossiping and broadcasting in communication networks. *Networks*, 18(4):319–349, 1988.
- [19] Wei Ye, John Heidemann, and Deborah Estrin. An energy-efficient mac protocol for wireless sensor networks. In *Proceedings of the IEEE Infocom*, pages 1567–1576, USA, June 2002.
- [20] Stephanie Lindsey and Cauligi Raghavendra. Pegasus: Power-efficient gathering in sensor information systems. In *IEEE Aerospace Conference Proceedings*, volume 3, pages 3–1125–3–1130, March 2002.
- [21] Wendi Rabiner Heinzelman, Anantha Chandrakasan, and Hari Balakrishnan. Energy-Efficient Communication Protocol for Wireless Microsensor Networks. In *Hawaii International Conference on System Sciences*, pages 3005–3014, January 2000.
- [22] Stanislava Soro and Wendi Heinzelman. Cluster head election techniques for coverage preservation in wireless sensor networks. *Ad Hoc Networks*, 7(5):955–972, July 2009.
- [23] Joanna Kulik, Wendi Heinzelman, and Hari Balakrishnan. Negotiation-based protocols for disseminating information in wireless sensor networks. *Wireless Networks*, 8:169–185, March 2002.
- [24] Ossama Younis and Sonia Fahmy. HEED: A hybrid, energy-efficient, distributed clustering approach for ad-hoc sensor networks. *IEEE Transactions on Mobile Computing*, 3(4):366–379, 2004.

- [25] Qi Dong and Donggang Liu. Resilient Cluster Leader Election for Wireless Sensor Networks. In *IEEE Communications Society Conference on Sensor, Mesh and Ad Hoc Communications and Networks*, pages 108–116, June 2009.
- [26] Si-Yu Ou, Hung-Chang Hsiao, Chi-Kuo Chiang, and Chung-Ta King. Designing power-aware overlays in heterogeneous wireless sensor networks. In *IEEE International Conference on Sensor Networks, Ubiquitous, and Trustworthy Computing*, pages 488–498, June 2006.
- [27] Georgios Smaragdakis, Ibrahim Matta, and Azer Bestavros. SEP: A stable election protocol for clustered heterogeneous wireless sensor networks. In *Second International Workshop on Sensor and Actor Network Protocols and Applications*, August 2004. Available at http://csr.bu.edu/sep/SEP_SANPA04.pdf.
- [28] Marcin Brzozowski, Hendrik Salomon, and Peter Langendoerfer. Completely Distributed Low Duty Cycle Communication for Long-Living Sensor Networks. *IEEE International Conference on Computational Science and Engineering*, 2:109–116, August 2009.
- [29] Wen Hu, Chun Tung Chou, Sanjay Jha, and Nirupama Bulusu. Deploying long-lived and cost-effective hybrid sensor networks. *Ad Hoc Networks*, 4(6):749–767, November 2006.
- [30] Vivek Mhatre and Catherine Rosenberg. Homogeneous vs heterogeneous clustered sensor networks: a comparative study. In *IEEE International Conference on Communications*, volume 6, pages 3646–3651, June 2004.

- [31] Mark Yarvis, Nandakishore Kushalnagar, Harkirat Singh, Anand Rangarajan, York Liu, and Suresh Singh. Exploiting heterogeneity in sensor networks. In *INFOCOM 2005. 24th Annual Joint Conference of the IEEE Computer and Communications Societies. Proceedings IEEE*, volume 2, pages 878–890, March 2005.
- [32] Andras Varga. OMNeT++. <http://www.omnetpp.org/>, Accessed on October 29, 2010.
- [33] National Information and Communications Technology Australia. Castalia. <http://castalia.npc.nicta.com.au/>, Accessed on January 19, 2012.
- [34] Pierre L'Ecuyer and Eric Buist. Simulation in java with ssj. In *Proceedings of the 37th conference on Winter simulation*, WSC '05, pages 611–620, December 2005.
- [35] Oracle Labs. Oracle SunSPOTs. <http://www.sunspotworld.com>, Accessed on January 25, 2011.
- [36] Oracle Labs. Java 2 Micro Edition. <http://www.oracle.com/technetwork/java/javame/overview/index.html>, Accessed on March 16, 2011.
- [37] Oracle Labs. Sun sunspot developers manual. <http://sunspotworld.com/docs/Yellow/SunSPOT-Programmers-Manual.pdf>, Accessed on March 16, 2011.
- [38] Seeed Studio. Seeeduino Stalker. <http://www.seeedstudio.com/depot/seeeduino-stalker-p-600.html>, Accessed on February 1, 2011.

- [39] Arduino. Arduino project. <http://www.arduino.cc/>, Accessed on March 16, 2011.
- [40] Andrew Rapp. Xbee-arduino. <http://code.google.com/p/xbee-arduino/>, Accessed on March 17, 2011.
- [41] Joseph Polastre, Robert Szewczyk, and David Culler. Telos: enabling ultra-low power wireless research. In *Information Processing in Sensor Networks. Fourth International Symposium on*, pages 364–369, April 2005.
- [42] Jun Luo, Jacques Panchard, Micha Pirkowski, Matthias Grossglauser, and Jean-Pierre Hubaux. MobiRoute: Routing towards a mobile sink for improving lifetime in sensor networks. In *Distributed Computing in Sensor Systems*, volume 4026 of *Lecture Notes in Computer Science*, pages 480–497. Springer Verlag, 2006.
- [43] Ronald Rivest, Adi Shamir, and Leonard Adleman. A method for obtaining digital signatures and public-key cryptosystems. *Commun. ACM*, 21(2):120–126, February 1978.
- [44] Eugene Shih, Seong-Hwan Cho, Nathan Ickes, Rex Min, Amit Sinha, Alice Wang, and Anantha Chandrakasan. Physical layer driven protocol and algorithm design for energy-efficient wireless sensor networks. In *Proceedings of the 7th annual international conference on Mobile computing and networking*, MobiCom, pages 272–287, 2001.
- [45] Shan Lin, Jingbin Zhang, Gang Zhou, Lin Gu, John A. Stankovic, and Tian He. ATPC: Adaptive transmission power control for wireless sensor networks. In

- Proceedings of the 4th international conference on Embedded networked sensor systems*, ACM Conference on Embedded Networked Sensor Systems, pages 223–236, 2006.
- [46] Liqiang Zheng, Wensi Wang, Alan Mathewson, Brendan O’Flynn, and Michael Hayes. An adaptive transmission power control method for wireless sensor networks. In *Signals and Systems Conference (ISSC 2010), IET Irish*, pages 261–265, June 2010.
- [47] Jian Xiao and Fengqi Yu. An efficient transmission power control algorithm in wireless sensor networks. In *Wireless Communications, Networking and Mobile Computing*, pages 2767–2770, September 2007.
- [48] Hsu Kahn, Victor Hsu, Joseph Kahn, and Kris Pister. Wireless communications for smart dust. In *Electronics Research Laboratory Technical Memorandum M98/2*, 1998.
- [49] Joseph Kahn, Randy Katz, and Kris Pister. Next century challenges: mobile networking for smart dust. In *Proceedings of the 5th annual ACM/IEEE international conference on Mobile computing and networking*, MobiCom, pages 271–278, 1999.
- [50] Jean-Philippe Vasseur and Adam Dunkels. *Interconnecting Smart Objects with IP - The Next Internet*. Morgan Kaufmann, 2010.
- [51] Adam Dunkels, Bjorn Gronvall, and Thiemo Voigt. Contiki - a lightweight and flexible operating system for tiny networked sensors. In *Proceedings of the First*

IEEE Workshop on Embedded Networked Sensors, pages 455 – 462, USA, November 2004.

- [52] Adam Dunkels. Contiki OS. <http://www.contiki-os.org/>, Accessed on March 1, 2012.

- [53] Adam Dunkels, Joakim Eriksson, and Nicolas Tsiftes. Low-power Interoperability for the IPv6-based Internet of Things. In *Proceedings of the 10th Scandinavian Workshop on Wireless Ad-hoc Networks*, Sweden, May 2011. <http://www.wirelessadhoc.org/index.php/program>.

- [54] Jonathan W. Hui and David E. Culler. Ipv6 in low-power wireless networks. *Proceedings of the IEEE*, 98(11):1865–1878, November 2010.

- [55] Steve Deering and Robert Hinden. Rfc 2460: Internet protocol, version 6 (ipv6) specification. <http://www.ietf.org/rfc/rfc2460.txt>, Accessed on March 1, 2012.